

UHViT: Complete Training Pipeline for IF 5+ Journal

Skin Lesion Classification with Uncertainty Quantification

#Install Dependencies (Fixed for RunPod)

Upgrade all packages to numpy 2.x compatible versions

```
!pip install -q --upgrade numpy==2.0.2
!pip install -q --upgrade matplotlib>=3.8.0 scipy>=1.12.0 scikit-learn>=1.4.0
!pip install -q --upgrade timm>=1.0.0
!pip install -q --upgrade wandb>=0.17.0 # Fixed: needs numpy 2.x compatible version
```

Now install the rest

```
!pip install -q albumentations>=1.3.1
!pip install -q kaggle==1.5.16
!pip install -q grad-cam
!pip install -q pandas>=2.0.3
!pip install -q seaborn>=0.13.0
!pip install -q tqdm>=4.66.1
!pip install -q ptflops>=0.7
!pip install -q gdown
!pip install -q opencv-python-headless
```

CELL 1B: EXTRA DEPENDENCY FOR EXTERNAL DATASET

```
!pip install -q medmnist
print("✓ medmnist installed")
```

✓ medmnist installed

CELL 2: IMPORTS

```
import os, sys, json, random, time, warnings
from pathlib import Path
from collections import Counter
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm.auto import tqdm
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from torch.optim import AdamW
```

```

from torch.optim.lr_scheduler import CosineAnnealingWarmRestarts
import timm
from PIL import Image
import albumentations as A
from albumentations.pytorch import ToTensorV2
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, balanced_accuracy_score,
f1_score, confusion_matrix, roc_auc_score,
precision_recall_fscore_support, roc_curve, auc
from sklearn.calibration import calibration_curve
from scipy import stats
import wandb
warnings.filterwarnings('ignore')
print(f"PyTorch: {torch.__version__}, CUDA:
{torch.cuda.is_available()}")
if torch.cuda.is_available(): print(f"GPU:
{torch.cuda.get_device_name(0)}")

```

PyTorch: 2.8.0+cu128, CUDA: True

GPU: NVIDIA GeForce RTX 5090

CELL 3: CONFIGURATION

```

CONFIG = {
    'KAGGLE_USERNAME': 'sushovanchaudhury', # UPDATE THIS
    'KAGGLE_KEY': '6ad1238ff7f05a1d17687883ac1cb80a', #
    'WANDB_API_KEY': '07ac9d1305afd5973d9022f150eb134c3b77ca7b',
    # UPDATE THIS
    'seed': 42, 'n_folds': 5, 'epochs': 30, 'batch_size': 32,
    'img_size': 224,
    'num_workers': 4, 'lr': 1e-4, 'weight_decay': 1e-4,
    'mc_dropout_samples': 10,
    'num_classes': 8, 'dropout_rate': 0.3,
    'swin_model': 'swin_tiny_patch4_window7_224',
    'efficientnet_model': 'efficientnet_b3',
    'class_names': ['AK', 'BCC', 'BKL', 'DF', 'MEL', 'NV', 'SCC',
    'VASC'],
    'data_dir': './data', 'output_dir': './outputs', 'checkpoint_dir':
    './checkpoints',
    'device': 'cuda' if torch.cuda.is_available() else 'cpu'
}
for d in [CONFIG['data_dir'], CONFIG['output_dir'],
CONFIG['checkpoint_dir']]:
    Path(d).mkdir(parents=True, exist_ok=True)
def set_seed(seed):
    random.seed(seed); np.random.seed(seed); torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed);
torch.backends.cudnn.deterministic = True
set_seed(CONFIG['seed'])

```

```
print(f"✓ Config set. Device: {CONFIG['device']}")
```

✓ Config set. Device: cuda

CELL 4: SETUP APIs

```
os.makedirs(os.path.expanduser('~/.kaggle'), exist_ok=True)
with open(os.path.expanduser('~/.kaggle/kaggle.json'), 'w') as f:
    json.dump({"username": CONFIG['KAGGLE_USERNAME'], "key":
CONFIG['KAGGLE_KEY']}, f)
os.chmod(os.path.expanduser('~/.kaggle/kaggle.json'), 0o600)
os.environ['WANDB_API_KEY'] = CONFIG['WANDB_API_KEY']
wandb.login(key=CONFIG['WANDB_API_KEY'])
print("✓ APIs configured")
```

wandb: WARNING If you're specifying your api key in code, ensure this code is not shared publicly.

wandb: WARNING Consider setting the WANDB_API_KEY environment variable, or running `wandb login` from the command line.

wandb: No netrc file found, creating one.

wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

wandb: Currently logged in as: sushovan-chaudhury (sushovan-chaudhury-bits-pilani) to https://api.wandb.ai. Use `wandb login --relogin` to force relogin

✓ APIs configured

```
!kaggle datasets list -s pad-ufes -p 5
```

Warning: Looks like you're using an outdated API Version, please consider updating (server 1.8.3 / client 1.5.16)

No datasets found

CELL 5: DOWNLOAD DATASETS (REPLACE THIS ENTIRE CELL)

```
from pathlib import Path
```

```
isic_path = Path(CONFIG['data_dir']) / 'isic-2019'
```

```
if not isic_path.exists():
```

```
    !kaggle datasets download -d andrewmvd/isic-2019 -p
{CONFIG['data_dir']}
```

```
    !unzip -q {CONFIG['data_dir']}/isic-2019.zip -d
```

```
{CONFIG['data_dir']}/isic-2019
```

```
    !rm {CONFIG['data_dir']}/isic-2019.zip
```

```
ham_path = Path(CONFIG['data_dir']) / 'ham10000'
```

```
if not ham_path.exists():
```

```
    !kaggle datasets download -d kmader/skin-cancer-mnist-ham10000 -p
{CONFIG['data_dir']}
```

```
    !unzip -q {CONFIG['data_dir']}/skin-cancer-mnist-ham10000.zip -d
```

```
{CONFIG['data_dir']}/ham10000
!rm {CONFIG['data_dir']}/skin-cancer-mnist-ham10000.zip
```

```
print("✓ Datasets ready (ISIC2019, HAM10000)")
```

```
# FIX: find correct ISIC2019 CSV + image directory (robust)
```

```
from pathlib import Path
import pandas as pd
```

```
isic_base = Path(CONFIG["data_dir"]) / "isic-2019"
```

```
# 1) Find a CSV that has an 'image' column and class columns
```

```
candidate csvs = list(isic_base.rglob("*.csv"))
```

```
best_csv = None
```

```
for p in candidate_csvs:
```

```
try:
```

```
df = pd.read_csv(p, nrows=5)
```

```
if "image" in df.columns:
```

```
best_csv = p
```

break

```
except:
```

pass

```
print("Chosen CSV:", best_csv)
```

```
train_df = pd.read_csv(best_csv)
```

```
# 2) Find an images directory that contains lots of .jpg files
```

```
candidate_dirs = [d for d in isic.base.rglob("*") if d.is_dir()]
```

```
def jpg_count(d):
```

```
return len(list(d.glob("*.jpg"))) + len(list(d.glob("*.JPG"))) +
```

```
len(list(d.glob("*.jpeg"))) + len(list(d.glob("*.png")))
```

```
best_img_dir = max(candidate_dirs, key=jpg_count)
```

```
print("Chosen image dir:", best_img_dir, "| image files:",
```

```
jpg count(best_img_dir))
```

```

train_img_dir = best_img_dir
print("train_df columns:", train_df.columns[:15])
print("train_df size:", len(train_df))

Chosen CSV: data/isic-2019/ISIC_2019_Training_GroundTruth.csv
Chosen image dir:
data/isic-2019/ISIC_2019_Training_Input/ISIC_2019_Training_Input |
image files: 25331
train_df columns: Index(['image', 'MEL', 'NV', 'BCC', 'AK', 'BKL',
'DF', 'VASC', 'SCC', 'UNK'], dtype='object')
train_df size: 25331

def get_train_transforms(img_size):
    return A.Compose([
        A.RandomResizedCrop(size=(img_size, img_size), scale=(0.8,
1.0)),
        A.HorizontalFlip(p=0.5),
        A.VerticalFlip(p=0.5),
        A.RandomRotate90(p=0.5),
        A.ShiftScaleRotate(shift_limit=0.1, scale_limit=0.1,
rotate_limit=45, p=0.5),
        A.OneOf([
            A.GaussNoise(var_limit=(10, 50)),
            A.GaussianBlur(blur_limit=(3, 7))
        ], p=0.3),
        A.OneOf([
            A.CLAHE(clip_limit=4.0),
            A.RandomBrightnessContrast()
        ], p=0.5),
        A.HueSaturationValue(p=0.5),
        A.CoarseDropout(
            num_holes_range=(1, 8),
            hole_height_range=(img_size//16, img_size//8),
            hole_width_range=(img_size//16, img_size//8),
            p=0.3
        ),
        A.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225]),
        ToTensorV2()
    ])

def get_val_transforms(img_size):
    return A.Compose([
        A.Resize(height=img_size, width=img_size),
        A.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225]),
        ToTensorV2()
    ])

```

```

import albumentations as A
import albumentations
print("Albumentations version:", albumentations.__version__)

# instantiate to force validation now
t = get_train_transforms(CONFIG["img_size"])
print("Train transforms OK:", t)

Albumentations version: 2.0.8
Train transforms OK: Compose([
  RandomResizedCrop(p=1.0, area_for_downscale=None, interpolation=1,
mask_interpolation=0, ratio=(0.75, 1.3333333333333333), scale=(0.8,
1.0), size=(224, 224)),
  HorizontalFlip(p=0.5),
  VerticalFlip(p=0.5),
  RandomRotate90(p=0.5),
  ShiftScaleRotate(p=0.5, shift_limit_x=(-0.1, 0.1), shift_limit_y=(-
0.1, 0.1), scale_limit=(-0.09999999999999998, 0.100000000000000009),
rotate_limit=(-45.0, 45.0), interpolation=1, border_mode=0, fill=0.0,
fill_mask=0.0, rotate_method='largest_box', mask_interpolation=0),
  OneOf([
    GaussNoise(p=0.5, mean_range=(0.0, 0.0), noise_scale_factor=1.0,
per_channel=True, std_range=(0.2, 0.44)),
    GaussianBlur(p=0.5, blur_limit=(3, 7), sigma_limit=(0.5, 3.0)),
  ], p=0.3),
  OneOf([
    CLAHE(p=0.5, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8)),
    RandomBrightnessContrast(p=0.5, brightness_by_max=True,
brightness_limit=(-0.2, 0.2), contrast_limit=(-0.2, 0.2),
ensure_safe_range=False),
  ], p=0.5),
  HueSaturationValue(p=0.5, hue_shift_limit=(-20.0, 20.0),
sat_shift_limit=(-30.0, 30.0), val_shift_limit=(-20.0, 20.0)),
  CoarseDropout(p=0.3, fill=0.0, fill_mask=None,
hole_height_range=(14, 28), hole_width_range=(14, 28),
num_holes_range=(1, 8)),
  Normalize(p=1.0, max_pixel_value=255.0, mean=(0.485, 0.456, 0.406),
normalization='standard', std=(0.229, 0.224, 0.225)),
  ToTensorV2(p=1.0, transpose_mask=False),
], p=1.0, bbox_params=None, keypoint_params=None,
additional_targets={}, is_check_shapes=True)

# CELL 7: DATASET CLASS
class ISIC2019Dataset(Dataset):
    def __init__(self, df, img_dir, transform=None):
        self.df = df.reset_index(drop=True)
        self.img_dir = Path(img_dir)
        self.transform = transform

    def __len__(self): return len(self.df)
    def __getitem__(self, idx):

```

```

        row = self.df.iloc[idx]
        img_name = row['image']
        img_path = None
        for ext in ['.jpg', '.jpeg', '.png', '.JPG']:
            p = self.img_dir / f"{img_name}{ext}"
            if p.exists(): img_path = p; break
        if not img_path:
            for f in self.img_dir.glob(f"{img_name}*"): img_path = f;
break
        image = np.array(Image.open(img_path).convert('RGB'))
        label = 0
        for i, cls in enumerate(CONFIG['class_names']):
            if cls in row and row[cls] == 1.0: label = i; break
        if self.transform: image = self.transform(image=image)
['image']
        return image, label

class HAM10000Dataset(Dataset):
    def __init__(self, csv_path, img_dirs, transform=None):
        self.df = pd.read_csv(csv_path)
        self.img_dirs = [Path(d) for d in img_dirs] if
isinstance(img_dirs, list) else [Path(img_dirs)]
        self.transform = transform
        self.label_map = {'akiec': 0, 'bcc': 1, 'bkl': 2, 'df': 3,
'mel': 4, 'nv': 5, 'vasc': 7}
    def __len__(self): return len(self.df)
    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = None
        for d in self.img_dirs:
            for ext in ['.jpg', '.jpeg', '.png']:
                p = d / f"{row['image_id']}{ext}"
                if p.exists(): img_path = p; break
            if img_path: break
        image = np.array(Image.open(img_path).convert('RGB'))
        label = self.label_map.get(row['dx'], 0)
        if self.transform: image = self.transform(image=image)
['image']
        return image, label
print("✓ Dataset classes defined")

```

✓ Dataset classes defined

CELL 8: LOAD DATA (REPLACE THIS ENTIRE CELL)

```

from pathlib import Path
import pandas as pd

```

```

isic_base = Path(CONFIG["data_dir"]) / "isic-2019"

```

```

# 1) Pick the correct LABEL CSV (GroundTruth), not the Metadata CSV
csv_candidates = list(isic_base.rglob("*.csv"))
gt_candidates = [p for p in csv_candidates if ("ground" in
p.name.lower()) or ("truth" in p.name.lower())]

if len(gt_candidates) == 0:
    raise FileNotFoundError(
        "GroundTruth CSV not found. Expected something like
'ISIC_2019_Training_GroundTruth.csv' "
        "inside data/isic-2019/. Please check your unzip output."
    )

# Prefer the one that contains 'Training' if multiple exist
gt_candidates_sorted = sorted(gt_candidates, key=lambda p: ("training"
not in p.name.lower(), len(p.name)))
train_csv = gt_candidates_sorted[0]
train_df = pd.read_csv(train_csv)

print("Using label CSV:", train_csv)
print("train_df size:", len(train_df))

# 2) Find the correct IMAGE directory (folder with most jpg/png)
candidate_dirs = [d for d in isic_base.rglob("*") if d.is_dir()]

def image_count(d: Path) -> int:
    return (
        len(list(d.glob("*.jpg")))
        + len(list(d.glob("*.JPG")))
        + len(list(d.glob("*.jpeg")))
        + len(list(d.glob("*.png")))
    )

train_img_dir = max(candidate_dirs, key=image_count)
print("Using image dir:", train_img_dir, "| image files:",
image_count(train_img_dir))

# 3) Verify the CSV actually contains the class columns
missing_cols = [c for c in CONFIG["class_names"] if c not in
train_df.columns]
if len(missing_cols) > 0:
    raise ValueError(
        f"This CSV does not contain required class columns:
{missing_cols}\n"
        f"You likely loaded the Metadata CSV by mistake.\n"
        f"Loaded columns (first 20): {list(train_df.columns[:20])}\n"
        f"CSV path used: {train_csv}"
    )

print("✓ Class columns found:", [c for c in CONFIG["class_names"] if c

```



```

in train_df.columns])
print("Class distribution (sum of one-hot columns):")
for cls in CONFIG["class_names"]:
    print(f" {cls}: {int(train_df[cls].sum())}")

Using label CSV: data/isic-2019/ISIC_2019_Training_GroundTruth.csv
train_df size: 25331
Using image dir:
data/isic-2019/ISIC_2019_Training_Input/ISIC_2019_Training_Input |
image files: 25331
✓ Class columns found: ['AK', 'BCC', 'BKL', 'DF', 'MEL', 'NV', 'SCC',
'VASC']
Class distribution (sum of one-hot columns):
AK: 867
BCC: 3323
BKL: 2624
DF: 239
MEL: 4522
NV: 12875
SCC: 628
VASC: 253

ds = ISIC2019Dataset(train_df.iloc[:200], train_img_dir,
get_train_transforms(CONFIG["img_size"]))
x, y = ds[0]
print("Sample tensor:", x.shape, "label:", y)
print("Unique labels in first 200:", sorted({ds[i][1] for i in
range(200)}))

Sample tensor: torch.Size([3, 224, 224]) label: 5
Unique labels in first 200: [4, 5]

# Build labels exactly like your dataset does
labels = []
for i in range(len(train_df)):
    row = train_df.iloc[i]
    found = False
    for j, cls in enumerate(CONFIG["class_names"]):
        if cls in train_df.columns and row[cls] == 1.0:
            labels.append(j)
            found = True
            break
    if not found:
        labels.append(-1)

labels = np.array(labels)
print("Label counts:", np.bincount(labels[labels>=0],
minlength=CONFIG["num_classes"]))
print("Missing labels (-1):", np.sum(labels == -1))
print("Unique labels:", np.unique(labels))

```

```
Label counts: [ 867 3323 2624 239 4522 12875 628 253]
Missing labels (-1): 0
Unique labels: [0 1 2 3 4 5 6 7]
```

```
# CELL 9: UHViT MODEL
```

```
class GatedFusion(nn.Module):
    def __init__(self, dim):
        super().__init__()
        self.gate = nn.Sequential(nn.Linear(dim * 2, dim),
nn.Sigmoid())
        self.proj = nn.Linear(dim * 2, dim)
    def forward(self, x1, x2):
        concat = torch.cat([x1, x2], dim=-1)
        return self.proj(concat) + self.gate(concat) * x1 + (1 -
self.gate(concat)) * x2

class MCDropout(nn.Module):
    def __init__(self, p=0.3): super().__init__(); self.p = p
    def forward(self, x): return F.dropout(x, p=self.p, training=True)

class UHViT(nn.Module):
    def __init__(self, num_classes=8, dropout_rate=0.3):
        super().__init__()
        self.swin = timm.create_model(CONFIG['swin_model'],
pretrained=True, num_classes=0)
        self.efficientnet =
timm.create_model(CONFIG['efficientnet_model'], pretrained=True,
num_classes=0)
        self.swin_proj = nn.Linear(self.swin.num_features, 512)
        self.eff_proj = nn.Linear(self.efficientnet.num_features, 512)
        self.fusion = GatedFusion(512)
        self.mc_dropout = MCDropout(dropout_rate)
        self.classifier = nn.Sequential(nn.LayerNorm(512),
nn.Linear(512, 256), nn.GELU(), MCDropout(dropout_rate),
nn.Linear(256, num_classes))
    def forward(self, x):
        swin_feat = self.swin_proj(self.swin(x))
        eff_feat = self.eff_proj(self.efficientnet(x))
        return self.classifier(self.mc_dropout(self.fusion(swin_feat,
eff_feat)))
model = UHViT(); print(f"✓ UHViT: {sum(p.numel() for p in
model.parameters()):,} params"); del model
```

```
{"model_id": "427402726956412c9efceb50f45333c4", "version_major": 2, "vers
ion_minor": 0}
```

```
{"model_id": "67db1b3300b5429b8e83f8a9584c0eda", "version_major": 2, "vers
ion_minor": 0}
```

✓ UHViT: 40,580,266 params

CELL 10: ABLATION & SOTA MODELS

```
class SwinOnly(nn.Module):
    def __init__(self, num_classes=8, dropout_rate=0.3):
        super().__init__()
        self.swin = timm.create_model(CONFIG['swin_model'],
pretrained=True, num_classes=0)
        self.classifier =
nn.Sequential(nn.LayerNorm(self.swin.num_features),
nn.Dropout(dropout_rate), nn.Linear(self.swin.num_features,
num_classes))
    def forward(self, x): return self.classifier(self.swin(x))

class EfficientNetOnly(nn.Module):
    def __init__(self, num_classes=8, dropout_rate=0.3):
        super().__init__()
        self.eff = timm.create_model(CONFIG['efficientnet_model'],
pretrained=True, num_classes=0)
        self.classifier =
nn.Sequential(nn.LayerNorm(self.eff.num_features),
nn.Dropout(dropout_rate), nn.Linear(self.eff.num_features,
num_classes))
    def forward(self, x): return self.classifier(self.eff(x))

class UHViTConcatFusion(nn.Module):
    def __init__(self, num_classes=8, dropout_rate=0.3):
        super().__init__()
        self.swin = timm.create_model(CONFIG['swin_model'],
pretrained=True, num_classes=0)
        self.efficientnet =
timm.create_model(CONFIG['efficientnet_model'], pretrained=True,
num_classes=0)
        self.swin_proj = nn.Linear(self.swin.num_features, 512)
        self.eff_proj = nn.Linear(self.efficientnet.num_features, 512)
        self.classifier = nn.Sequential(nn.LayerNorm(1024),
nn.Dropout(dropout_rate), nn.Linear(1024, num_classes))
    def forward(self, x): return
self.classifier(torch.cat([self.swin_proj(self.swin(x)),
self.eff_proj(self.efficientnet(x))], dim=-1))

SOTA_MODELS = ['resnet50', 'efficientnet_b3', 'vit_base_patch16_224',
'swin_tiny_patch4_window7_224', 'convnext_tiny', 'densenet121']
def create_sota_model(name, num_classes=8): return
timm.create_model(name, pretrained=True, num_classes=num_classes)
print("✓ Ablation & SOTA models defined")
```

✓ Ablation & SOTA models defined

CELL 11: LOSS FUNCTIONS

```
class FocalLoss(nn.Module):
    def __init__(self, alpha=None, gamma=2.0): super().__init__();
    self.alpha = alpha; self.gamma = gamma
    def forward(self, inputs, targets):
        ce = F.cross_entropy(inputs, targets, weight=self.alpha,
reduction='none')
        return (((1 - torch.exp(-ce)) ** self.gamma) * ce).mean()

class ClassBalancedLoss(nn.Module):
    def __init__(self, samples_per_class, num_classes, beta=0.9999,
gamma=2.0):
        super().__init__()
        effective_num = 1.0 - np.power(beta, samples_per_class)
        weights = (1.0 - beta) / np.array(effective_num)
        self.weights = torch.tensor(weights / np.sum(weights) *
num_classes, dtype=torch.float32)
        self.gamma = gamma
    def forward(self, inputs, targets):
        self.weights = self.weights.to(inputs.device)
        ce = F.cross_entropy(inputs, targets, weight=self.weights,
reduction='none')
        return (((1 - torch.exp(-ce)) ** self.gamma) * ce).mean()

def get_class_weights(labels):
    counts = np.maximum(np.bincount(labels,
minlength=CONFIG['num_classes']), 1)
    weights = 1.0 / counts
    return torch.tensor(weights / weights.sum() * len(weights),
dtype=torch.float32)

def create_balanced_sampler(labels):
    counts = np.maximum(np.bincount(labels,
minlength=CONFIG['num_classes']), 1)
    sample_weights = (1.0 / counts)[labels]
    return WeightedRandomSampler(sample_weights, len(labels),
replacement=True)
print("✓ Loss functions defined")
```

✓ Loss functions defined

CELL 12: TRAINING FUNCTIONS

```
def train_epoch(model, loader, criterion, optimizer, scheduler,
device, epoch):
    model.train()
    running_loss, all_preds, all_labels = 0.0, [], []
    pbar = tqdm(loader, desc=f'Train E{epoch}')
    for images, labels in pbar:
        images, labels = images.to(device), labels.to(device)
```

```

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(),
max_norm=1.0)
        optimizer.step()
        if scheduler: scheduler.step()
        running_loss += loss.item()
        all_preds.extend(outputs.argmax(dim=1).cpu().numpy())
        all_labels.extend(labels.cpu().numpy())
        pbar.set_postfix({'loss': f'{loss.item():.4f}'})
    return running_loss / len(loader), accuracy_score(all_labels,
all_preds), balanced_accuracy_score(all_labels, all_preds)

def validate(model, loader, criterion, device):
    model.eval()
    running_loss, all_preds, all_labels, all_probs = 0.0, [], [], []
    with torch.no_grad():
        for images, labels in tqdm(loader, desc='Val'):
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            running_loss += criterion(outputs, labels).item()
            all_probs.extend(F.softmax(outputs, dim=1).cpu().numpy())
            all_preds.extend(outputs.argmax(dim=1).cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
    all_preds, all_labels, all_probs = np.array(all_preds),
np.array(all_labels), np.array(all_probs)
    return running_loss / len(loader), accuracy_score(all_labels,
all_preds), balanced_accuracy_score(all_labels, all_preds),
f1_score(all_labels, all_preds, average='macro'), all_preds,
all_labels, all_probs

def validate_with_tta(model, loader, device):
    model.eval()
    all_probs, all_labels = [], []
    with torch.no_grad():
        for images, labels in tqdm(loader, desc='TTA'):
            images = images.to(device)
            probs = torch.stack([F.softmax(model(images), dim=1),
F.softmax(model(torch.flip(images, [3])), dim=1),
F.softmax(model(torch.flip(images, [2])), dim=1)]).mean(0)
            all_probs.extend(probs.cpu().numpy());
    all_labels.extend(labels.numpy())
    all_probs, all_labels = np.array(all_probs), np.array(all_labels)
    all_preds = all_probs.argmax(axis=1)
    return accuracy_score(all_labels, all_preds),
balanced_accuracy_score(all_labels, all_preds), f1_score(all_labels,
all_preds, average='macro'), all_preds, all_labels, all_probs

```

```
print("✓ Training functions defined")
```

✓ Training functions defined

```
# CELL 13: UNCERTAINTY ESTIMATION
```

```
def estimate_uncertainty(model, loader, device, n_samples=10):
    model.train()
    all_means, all_uncertainties, all_labels = [], [], []
    with torch.no_grad():
        for images, labels in tqdm(loader, desc='Uncertainty'):
            images = images.to(device)
            outputs = torch.stack([F.softmax(model(images), dim=1) for
_ in range(n_samples)])
            all_means.extend(outputs.mean(0).cpu().numpy())

    all_uncertainties.extend(outputs.std(0).mean(1).cpu().numpy())
    all_labels.extend(labels.numpy())
    return np.array(all_means), np.array(all_uncertainties),
np.array(all_labels)

def plot_uncertainty_analysis(uncertainties, predictions, labels,
save_path):
    correct = (predictions == labels)
    fig, axes = plt.subplots(1, 3, figsize=(15, 4))
    axes[0].hist(uncertainties[correct], bins=30, alpha=0.7,
label='Correct', density=True)
    axes[0].hist(uncertainties[~correct], bins=30, alpha=0.7,
label='Incorrect', density=True)
    axes[0].legend(); axes[0].set_title('Uncertainty Distribution')
    thresholds = np.percentile(uncertainties, np.arange(0, 100, 5))
    accs, coverages = [], []
    for t in thresholds:
        mask = uncertainties <= t
        if mask.sum() > 0: accs.append(accuracy_score(labels[mask],
predictions[mask])); coverages.append(mask.mean())
    axes[1].plot(coverages, accs, 'o-');
    axes[1].set_xlabel('Coverage'); axes[1].set_ylabel('Accuracy');
    axes[1].set_title('Accuracy vs Coverage')
    class_unc = [uncertainties[labels == i].mean() if (labels ==
i).sum() > 0 else 0 for i in range(CONFIG['num_classes'])]
    axes[2].bar(CONFIG['class_names'], class_unc);
    axes[2].set_title('Per-Class Uncertainty');
    axes[2].tick_params(axis='x', rotation=45)
    plt.tight_layout(); plt.savefig(save_path, dpi=150);
    wandb.log({"uncertainty": wandb.Image(save_path)}); plt.close()
print("✓ Uncertainty functions defined")
```

✓ Uncertainty functions defined

CELL 14: VISUALIZATION FUNCTIONS

```
def plot_confusion_matrix(y_true, y_pred, class_names, save_path):
    cm = confusion_matrix(y_true, y_pred)
    cm_norm = cm.astype('float') / cm.sum(axis=1, keepdims=True)
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names, ax=axes[0])
    axes[0].set_title('Confusion Matrix');
    axes[0].set_xlabel('Predicted'); axes[0].set_ylabel('True')
    sns.heatmap(cm_norm, annot=True, fmt='.2f', cmap='Blues',
xticklabels=class_names, yticklabels=class_names, ax=axes[1])
    axes[1].set_title('Normalized'); axes[1].set_xlabel('Predicted');
    axes[1].set_ylabel('True')
    plt.tight_layout(); plt.savefig(save_path, dpi=150);
wandb.log({"confusion_matrix": wandb.Image(save_path)}); plt.close()

def plot_roc_curves(y_true, y_probs, class_names, save_path):
    fig, ax = plt.subplots(figsize=(10, 8))
    for i, cls in enumerate(class_names):
        binary = (y_true == i).astype(int)
        if 0 < binary.sum() < len(binary):
            fpr, tpr, _ = roc_curve(binary, y_probs[:, i])
            ax.plot(fpr, tpr, label=f'{cls} (AUC={auc(fpr,
tpr):.3f})')
    ax.plot([0, 1], [0, 1], 'k--'); ax.legend(loc='lower right')
    plt.savefig(save_path, dpi=150); wandb.log({"roc_curves":
wandb.Image(save_path)}); plt.close()

def plot_training_history(history, save_path):
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))
    epochs = range(1, len(history['train_loss']) + 1)
    axes[0, 0].plot(epochs, history['train_loss'], 'b-',
label='Train'); axes[0, 0].plot(epochs, history['val_loss'], 'r-',
label='Val'); axes[0, 0].set_title('Loss'); axes[0, 0].legend()
    axes[0, 1].plot(epochs, history['train_acc'], 'b-',
label='Train'); axes[0, 1].plot(epochs, history['val_acc'], 'r-',
label='Val'); axes[0, 1].set_title('Accuracy'); axes[0, 1].legend()
    axes[1, 0].plot(epochs, history['train_bacc'], 'b-',
label='Train'); axes[1, 0].plot(epochs, history['val_bacc'], 'r-',
label='Val'); axes[1, 0].set_title('Balanced Accuracy'); axes[1,
0].legend()
    axes[1, 1].plot(epochs, history['val_f1'], 'g-'); axes[1,
1].set_title('Macro F1')
    plt.tight_layout(); plt.savefig(save_path, dpi=150);
wandb.log({"training_history": wandb.Image(save_path)}); plt.close()
print("✓ Visualization functions defined")
```

✓ Visualization functions defined

CELL 15: CLINICAL & CALIBRATION ANALYSIS

```
def compute_ece(probs, labels, n_bins=15):
    bin_boundaries = np.linspace(0, 1, n_bins + 1)
    ece = 0.0
    for i in range(n_bins):
        mask = (probs > bin_boundaries[i]) & (probs <=
bin_boundaries[i + 1])
        if mask.sum() > 0: ece += mask.sum() * abs(probs[mask].mean()
- labels[mask].mean())
    return ece / len(probs)

def per_class_analysis(y_true, y_pred, y_probs, class_names):
    precision, recall, f1, support =
precision_recall_fscore_support(y_true, y_pred, average=None,
zero_division=0)
    auc_scores = []
    for i in range(len(class_names)):
        binary = (y_true == i).astype(int)
        try: auc_scores.append(roc_auc_score(binary, y_probs[:, i]) if
0 < binary.sum() < len(binary) else 0)
        except: auc_scores.append(0)
    df = pd.DataFrame({'Class': class_names, 'Precision': precision,
'Recall': recall, 'F1': f1, 'AUC': auc_scores, 'Support': support})
    wandb.log({"per_class_metrics": wandb.Table(dataframe=df)})
    return df

def plot_per_class_performance(df, save_path):
    fig, ax = plt.subplots(figsize=(14, 6))
    x = np.arange(len(df)); width = 0.2
    ax.bar(x - 1.5*width, df['Precision'], width, label='Precision');
ax.bar(x - 0.5*width, df['Recall'], width, label='Recall')
    ax.bar(x + 0.5*width, df['F1'], width, label='F1'); ax.bar(x +
1.5*width, df['AUC'], width, label='AUC')
    ax.set_xticks(x); ax.set_xticklabels(df['Class'], rotation=45);
ax.legend(); ax.set_ylim(0, 1)
    plt.tight_layout(); plt.savefig(save_path, dpi=150);
wandb.log({"per_class_chart": wandb.Image(save_path)}); plt.close()
print("✓ Clinical analysis functions defined")
```

✓ Clinical analysis functions defined

CELL 16: GRAD-CAM++

```
from pytorch_grad_cam import GradCAMPlusPlus
from pytorch_grad_cam.utils.image import show_cam_on_image
from pytorch_grad_cam.utils.model_targets import
ClassifierOutputTarget
```

```
def get_gradcam_visualizations(model, loader, device,
save_path='gradcam.png'):
```



```

model.eval()
try: target_layer = [model.efficientnet.conv_head]
except: target_layer = [list(model.efficientnet.children())[-2]]
cam = GradCAMPlusPlus(model=model, target_layers=target_layer)
fig, axes = plt.subplots(2, 4, figsize=(16, 8)); axes =
axes.flatten()
images, labels = next(iter(loader)); images =
images[:8].to(device); labels = labels[:8]
for i in range(min(8, len(images))):
    input_tensor = images[i:i + 1]
    with torch.no_grad(): pred_class =
model(input_tensor).argmax(dim=1).item()
    grayscale_cam = cam(input_tensor=input_tensor,
targets=[ClassifierOutputTarget(pred_class)][0, :])
    img = images[i].cpu().numpy().transpose(1, 2, 0)
    img = np.clip(img * np.array([0.229, 0.224, 0.225]) +
np.array([0.485, 0.456, 0.406]), 0, 1)
    vis = show_cam_on_image(img, grayscale_cam, use_rgb=True)
    axes[i].imshow(vis); axes[i].set_title(f'T:
{CONFIG["class_names"][labels[i]]} P:{CONFIG["class_names"]
[pred_class]}'); axes[i].axis('off')
plt.tight_layout(); plt.savefig(save_path, dpi=150);
wandb.log({"gradcam": wandb.Image(save_path)}); plt.close(); del cam
print("✓ Grad-CAM++ defined")

```

✓ Grad-CAM++ defined

CELL 17: STATISTICAL TESTS

```

def mcnemar_test(y_true, y_pred1, y_pred2):
    c1, c2 = (y_pred1 == y_true), (y_pred2 == y_true)
    b, c = np.sum(c1 & ~c2), np.sum(~c1 & c2)
    if b + c > 0: stat = (abs(b - c) - 1) ** 2 / (b + c); p = 1 -
stats.chi2.cdf(stat, df=1)
    else: stat, p = 0, 1.0
    return stat, p

def paired_t_test(scores1, scores2): return stats.ttest_rel(scores1,
scores2)
def wilcoxon_test(scores1, scores2):
    try: return stats.wilcoxon(scores1, scores2)
    except: return 0, 1.0

def compute_ci(scores, confidence=0.95):
    n, mean, se = len(scores), np.mean(scores), stats.sem(scores)
    h = se * stats.t.ppf((1 + confidence) / 2, n - 1)
    return mean, mean - h, mean + h
print("✓ Statistical tests defined")

```

✓ Statistical tests defined

CELL 18: COMPUTATIONAL ANALYSIS

```
from ptflops import get_model_complexity_info

def compute_model_complexity(model, input_size=(3, 224, 224)):
    try: macs, params = get_model_complexity_info(model, input_size,
as_strings=False, print_per_layer_stat=False, verbose=False); return
macs * 2, params
    except: return 0, sum(p.numel() for p in model.parameters())

def measure_inference_time(model, device, n_runs=100):
    model.eval(); dummy = torch.randn(1, 3, 224, 224).to(device)
    with torch.no_grad():
        for _ in range(10): _ = model(dummy)
    if device == 'cuda': torch.cuda.synchronize()
    times = []
    with torch.no_grad():
        for _ in range(n_runs):
            start = time.time(); _ = model(dummy)
            if device == 'cuda': torch.cuda.synchronize()
            times.append(time.time() - start)
    return np.mean(times) * 1000, np.std(times) * 1000

def computational_analysis(models_dict, device):
    results = []
    for name, model_fn in models_dict.items():
        print(f"Analyzing {name}..."); model = model_fn().to(device)
        flops, params = compute_model_complexity(model); mean_time,
std_time = measure_inference_time(model, device)
        results.append({'Model': name, 'Params (M)': params / 1e6,
'FLOPs (G)': flops / 1e9, 'Time (ms)': mean_time})
        del model; torch.cuda.empty_cache()
    df = pd.DataFrame(results); wandb.log({"computational":
wandb.Table(dataframe=df)}); return df
print("✓ Computational analysis defined")
```

✓ Computational analysis defined

CELL 19: MAIN TRAINING LOOP

```
def train_fold(fold, train_df, val_df, train_img_dir, device):
    print(f"\n{'='*50}\nFOLD {fold+1}/{CONFIG['n_folds']}\n{'='*50}")
    train_dataset = ISIC2019Dataset(train_df, train_img_dir,
get_train_transforms(CONFIG['img_size']))
    val_dataset = ISIC2019Dataset(val_df, train_img_dir,
get_val_transforms(CONFIG['img_size']))
    train_labels = []
    for i in range(len(train_df)):
        row = train_df.iloc[i]
```

```

        for j, cls in enumerate(CONFIG['class_names']):
            if cls in row and row[cls] == 1.0: train_labels.append(j);
break
        else: train_labels.append(0)
        train_labels = np.array(train_labels)
        sampler = create_balanced_sampler(train_labels)
        train_loader = DataLoader(train_dataset,
batch_size=CONFIG['batch_size'], sampler=sampler,
num_workers=CONFIG['num_workers'], pin_memory=True)
        val_loader = DataLoader(val_dataset,
batch_size=CONFIG['batch_size'], shuffle=False,
num_workers=CONFIG['num_workers'], pin_memory=True)
        model = UHViT(num_classes=CONFIG['num_classes'],
dropout_rate=CONFIG['dropout_rate']).to(device)
        samples_per_class = np.maximum(np.bincount(train_labels,
minlength=CONFIG['num_classes']), 1)
        criterion = ClassBalancedLoss(samples_per_class,
CONFIG['num_classes'])
        optimizer = AdamW(model.parameters(), lr=CONFIG['lr'],
weight_decay=CONFIG['weight_decay'])
        scheduler = CosineAnnealingWarmRestarts(optimizer, T_0=10,
T_mult=2)
        history = {'train_loss': [], 'train_acc': [], 'train_bacc': [],
'val_loss': [], 'val_acc': [], 'val_bacc': [], 'val_f1': []}
        best_bacc, best_state = 0, None
        for epoch in range(CONFIG['epochs']):
            train_loss, train_acc, train_bacc = train_epoch(model,
train_loader, criterion, optimizer, scheduler, device, epoch + 1)
            val_loss, val_acc, val_bacc, val_f1, val_preds, val_labels,
val_probs = validate(model, val_loader, criterion, device)
            wandb.log({'fold{fold}/train_loss': train_loss,
f'fold{fold}/train_bacc': train_bacc, f'fold{fold}/val_loss':
val_loss, f'fold{fold}/val_bacc': val_bacc, f'fold{fold}/val_f1':
val_f1})
            history['train_loss'].append(train_loss);
            history['train_acc'].append(train_acc);
            history['train_bacc'].append(train_bacc)
            history['val_loss'].append(val_loss);
            history['val_acc'].append(val_acc);
            history['val_bacc'].append(val_bacc); history['val_f1'].append(val_f1)
            print(f"E{epoch+1}: TrBAcc={train_bacc:.4f},
ValBAcc={val_bacc:.4f}, ValF1={val_f1:.4f}")
            if val_bacc > best_bacc: best_bacc = val_bacc; best_state =
model.state_dict().copy()
            model.load_state_dict(best_state); torch.save(best_state,
f"{CONFIG['checkpoint_dir']}/fold{fold}_best.pt")
            tta_acc, tta_bacc, tta_f1, tta_preds, tta_labels, tta_probs =
validate_with_tta(model, val_loader, device)
            print(f"TTA Results: Acc={tta_acc:.4f}, BAcc={tta_bacc:.4f},

```

```

f1={tta_f1:.4f}")
    plot_confusion_matrix(tta_labels, tta_preds,
CONFIG['class_names'], f"{CONFIG['output_dir']}/fold{fold}_cm.png")
    plot_roc_curves(tta_labels, tta_probs, CONFIG['class_names'],
f"{CONFIG['output_dir']}/fold{fold}_roc.png")
    plot_training_history(history,
f"{CONFIG['output_dir']}/fold{fold}_history.png")
    mean_probs, uncertainties, unc_labels =
estimate_uncertainty(model, val_loader, device)
    plot_uncertainty_analysis(uncertainties,
mean_probs.argmax(axis=1), unc_labels,
f"{CONFIG['output_dir']}/fold{fold}_uncertainty.png")
    get_gradcam_visualizations(model, val_loader, device,
f"{CONFIG['output_dir']}/fold{fold}_gradcam.png")
    pc_df = per_class_analysis(tta_labels, tta_preds, tta_probs,
CONFIG['class_names'])
    plot_per_class_performance(pc_df,
f"{CONFIG['output_dir']}/fold{fold}_perclass.png")
    return {'acc': tta_acc, 'bacc': tta_bacc, 'f1': tta_f1, 'preds':
tta_preds, 'labels': tta_labels, 'probs': tta_probs, 'model': model}
print("✓ Main training function defined")

```

✓ Main training function defined

```

# Sanity check: confirm first 50 images exist
missing = 0
for i in range(min(50, len(train_df))):
    img_name = str(train_df.iloc[i]["image"])
    found = False
    for ext in [".jpg", ".jpeg", ".png", ".JPG"]:
        if (Path(train_img_dir) / f"{img_name}{ext}").exists():
            found = True
            break
    if not found:
        # try glob
        if not list(Path(train_img_dir).glob(f"{img_name}*")):
            missing += 1
            if missing <= 5:
                print("Missing example:", img_name)

print("Missing in first 50:", missing)
assert missing == 0, "Image paths are wrong – fix train_img_dir / CSV
selection."

```

Missing in first 50: 0

```

# CELL 20: RUN 5-FOLD CROSS-VALIDATION
wandb.init(project='UHViT-SkinLesion-IF5', name='5fold-cv',
config=CONFIG)

```

```

all_labels = []
for i in range(len(train_df)):
    row = train_df.iloc[i]
    for j, cls in enumerate(CONFIG['class_names']):
        if cls in row and row[cls] == 1.0: all_labels.append(j); break
    else: all_labels.append(0)
all_labels = np.array(all_labels)
skf = StratifiedKFold(n_splits=CONFIG['n_folds'], shuffle=True,
random_state=CONFIG['seed'])
fold_results = []
for fold, (train_idx, val_idx) in enumerate(skf.split(train_df,
all_labels)):
    result = train_fold(fold, train_df.iloc[train_idx],
train_df.iloc[val_idx], train_img_dir, CONFIG['device'])
    fold_results.append(result)
    print(f"\nFold {fold+1} Complete: BAcc={result['bacc']:.4f},
F1={result['f1']:.4f}")
accs, baccs, f1s = [r['acc'] for r in fold_results], [r['bacc'] for r
in fold_results], [r['f1'] for r in fold_results]
print(f"\n{'='*50}\n5-FOLD CV RESULTS\n{'='*50}")
print(f"Accuracy: {np.mean(accs):.4f} ± {np.std(accs):.4f}")
print(f"Balanced Accuracy: {np.mean(baccs):.4f} ±
{np.std(baccs):.4f}")
print(f"Macro F1: {np.mean(f1s):.4f} ± {np.std(f1s):.4f}")
acc_mean, acc_lo, acc_hi = compute_ci(accs); bacc_mean, bacc_lo,
bacc_hi = compute_ci(baccs); f1_mean, f1_lo, f1_hi = compute_ci(f1s)
print(f"\n95% CI - Accuracy: [{acc_lo:.4f}, {acc_hi:.4f}]")
print(f"95% CI - Balanced Acc: [{bacc_lo:.4f}, {bacc_hi:.4f}]")
print(f"95% CI - Macro F1: [{f1_lo:.4f}, {f1_hi:.4f}]")
wandb.log({'final/acc_mean': np.mean(accs), 'final/acc_std':
np.std(accs), 'final/bacc_mean': np.mean(baccs), 'final/bacc_std':
np.std(baccs), 'final/f1_mean': np.mean(f1s), 'final/f1_std':
np.std(f1s)})

```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```

=====
FOLD 1/5
=====

```

```
{"model_id": "923139d0c37d46fbb263ccd4359d5688", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9594513e66b24825b20b6bf2fc3c8048", "version_major": 2, "version_minor": 0}
```

E1: TrBAcc=0.3389, ValBAcc=0.3773, ValF1=0.1059

```
{"model_id": "e422df2ca32e4bf5ae7e4a3d0df06b14", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7ff3fe94ce0a4c1db45f0454df583d58", "version_major": 2, "version_minor": 0}
```

E2: TrBAcc=0.4332, ValBAcc=0.4993, ValF1=0.1712

```
{"model_id": "32bcf77cc482484cb95240c18c18f538", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "130d6df8a6114523bd983445fde1c6ff", "version_major": 2, "version_minor": 0}
```

E3: TrBAcc=0.4644, ValBAcc=0.5360, ValF1=0.2290

```
{"model_id": "70fd35a31ddb47919727f8fd784a0648", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3d079d383d5b4d1ab7ea2e8ffc256f06", "version_major": 2, "version_minor": 0}
```

E4: TrBAcc=0.5364, ValBAcc=0.5593, ValF1=0.2700

```
{"model_id": "f17203636ea14ece8c7af6b599623d96", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "522a2d488747479b822c1369a546b3fd", "version_major": 2, "version_minor": 0}
```

E5: TrBAcc=0.5133, ValBAcc=0.5636, ValF1=0.2695

```
{"model_id": "e5eebf4767414750a09f37a8ed990dee", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "93f7e1df6bea43b882693b4d706b4e37", "version_major": 2, "version_minor": 0}
```

E6: TrBAcc=0.5456, ValBAcc=0.5989, ValF1=0.3447

```
{"model_id": "89b68b75dfe548c8adac679ef0f96117", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "53656a612f4c4a6db81fbd277aab2aa6", "version_major": 2, "version_minor": 0}
```

E7: TrBAcc=0.5959, ValBAcc=0.6240, ValF1=0.3721

```
{"model_id": "3f77d5167d6242ef94c3d7860198f8ae", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e1bfa5f4ef4b49e191e6875e3ac71e2d", "version_major": 2, "version_minor": 0}
```

E8: TrBAcc=0.6196, ValBAcc=0.6227, ValF1=0.3664

```
{"model_id": "ad9efae93da344b9987f3bbe145d5903", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4c56f56fcf36402b8eb4434a3991b5ec", "version_major": 2, "version_minor": 0}
```

E9: TrBAcc=0.5782, ValBAcc=0.5933, ValF1=0.3099

```
{"model_id": "37701483ca20472080f5168980c89025", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d89098d1ea9049e69529450869217c6a", "version_major": 2, "version_minor": 0}
```

E10: TrBAcc=0.5838, ValBAcc=0.6179, ValF1=0.3202

```
{"model_id": "54cc347cb7f44223886435f1bf410adb", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c5a377e04a66412fbca7c5f646ce5733", "version_major": 2, "version_minor": 0}
```

E11: TrBAcc=0.5989, ValBAcc=0.6425, ValF1=0.3897

```
{"model_id": "ada395667ef646428f93c056008a6665", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1402a1684c574ebb8d2c91791897e286", "version_major": 2, "version_minor": 0}
```

E12: TrBAcc=0.6274, ValBAcc=0.6358, ValF1=0.3766

```
{"model_id": "bda31714fe5f4deca071475239a4184d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9c1fa32577934ccd9d895a80f872600a", "version_major": 2, "version_minor": 0}
```

E13: TrBAcc=0.6501, ValBAcc=0.6479, ValF1=0.4311

```
{"model_id": "10bf40d09a194f0b840097bb2f511aa6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9eb1a1e885a44003bf4c90452f02171d", "version_major": 2, "version_minor": 0}
```

E14: TrBAcc=0.6765, ValBAcc=0.6679, ValF1=0.4526

```
{"model_id": "b38ebcae90a04fd7ac5756bbf760ad35", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "504884e197d24df99bc2e71f43934f1c", "version_major": 2, "version_minor": 0}
```

E15: TrBAcc=0.6858, ValBAcc=0.6786, ValF1=0.4672

```
{"model_id": "cee9ce3283d74b31b4e0429343fa7e20", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fe742e7843c746449a0d133ca944e13c", "version_major": 2, "version_minor": 0}
```

E16: TrBAcc=0.6939, ValBAcc=0.6781, ValF1=0.4700

```
{"model_id": "1a3f36ce991c4f349180e1b7806b2ec5", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "cdd40754901d4a088732afc0a65ad9a4", "version_major": 2, "version_minor": 0}
```

E17: TrBAcc=0.6385, ValBAcc=0.6240, ValF1=0.3254

```
{"model_id": "233865a080af4f7b8a482b40958aef77", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "16cc1171fa1e42928393aed500b5fb49", "version_major": 2, "version_minor": 0}
```

E18: TrBAcc=0.6211, ValBAcc=0.6576, ValF1=0.4019

```
{"model_id": "b3dee87e784c48d19ec68cb384f8bbb5", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b249c9ba52ee47d18d8e60ba0069e6fb", "version_major": 2, "version_minor": 0}
```

E19: TrBAcc=0.6381, ValBAcc=0.6549, ValF1=0.4266

```
{"model_id": "e59b1c9ca480457186246f370515f9c7", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "470c656d9e0b418e903a173e0fa481fe", "version_major": 2, "version_minor": 0}
```

E20: TrBAcc=0.6479, ValBAcc=0.6318, ValF1=0.3862


```
{"model_id": "cac0131e733d42a386b94d4f1875eb98", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "05c8d6da0fd1445e80ae47a774f3a4d0", "version_major": 2, "version_minor": 0}
```

E21: TrBAcc=0.6636, ValBAcc=0.6490, ValF1=0.3911

```
{"model_id": "499ba327e4b14c5ca1b743cf6004a768", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "06d7c2e086b1426b90bebbe499d1a8c5", "version_major": 2, "version_minor": 0}
```

E22: TrBAcc=0.6732, ValBAcc=0.6642, ValF1=0.4442

```
{"model_id": "970017e12ce2466ebfd85a3da2c613c2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "cef12ba1e0d346aa9d7125d9ed045b78", "version_major": 2, "version_minor": 0}
```

E23: TrBAcc=0.6833, ValBAcc=0.6808, ValF1=0.4572

```
{"model_id": "e293406bf5774f6ea9a97bd006fe1605", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b3859b3af32047c3a1cd2f8266f7cc01", "version_major": 2, "version_minor": 0}
```

E24: TrBAcc=0.7043, ValBAcc=0.6962, ValF1=0.4987

```
{"model_id": "0ddc57681ab6497db7c924150e748f2e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3dae5042de744f01b9cd44f43cfab331", "version_major": 2, "version_minor": 0}
```

E25: TrBAcc=0.7139, ValBAcc=0.7038, ValF1=0.5108

```
{"model_id": "227ea6ee3e694182b61d5e4ea9644635", "version_major": 2, "version_minor": 0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

```

    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "baa92eb34e3541048652ca7d59f9fc69", "version_major": 2, "vers
ion_minor": 0}

E26: TrBAcc=0.7206, ValBAcc=0.7080, ValF1=0.5501

{"model_id": "396ff0cf42fb4ce88b8a25c064e68fa7", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "9def7be33edb43bc8a4fd7cef7948d92", "version_major": 2, "vers
ion_minor": 0}

E27: TrBAcc=0.7316, ValBAcc=0.7007, ValF1=0.5318

{"model_id": "9aa8036b2b9d44ab85ad8d2a0df9a2f4", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "eac2c75ef48044f1884155d88de1bf13", "version_major": 2, "vers
ion_minor": 0}

E28: TrBAcc=0.7460, ValBAcc=0.7114, ValF1=0.5208

{"model_id": "1e0cb3a7f6644bad9f4c1f706f763ddf", "version_major": 2, "vers
ion_minor": 0}

```

```
{"model_id": "5c8c098199e342f5a7d351d77519842a", "version_major": 2, "version_minor": 0}
```

E29: TrBAcc=0.7531, ValBAcc=0.7082, ValF1=0.5439

```
{"model_id": "5b6d4694b35a45e1ba83d79c1da8d4b2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "58a5bf8b041243c280995d10a69ea38c", "version_major": 2, "version_minor": 0}
```

E30: TrBAcc=0.7605, ValBAcc=0.6981, ValF1=0.5602

```
{"model_id": "ef6fcf14b6584177a19d90bff0d35298", "version_major": 2, "version_minor": 0}
```

TTA Results: Acc=0.4796, BAcc=0.7173, F1=0.5722

```
{"model_id": "fd6aeb6e1d3e4b34a1ccc445ce27ab22", "version_major": 2, "version_minor": 0}
```

Fold 1 Complete: BAcc=0.7173, F1=0.5722

```
=====
FOLD 2/5
=====
```

```
{"model_id": "4c6bec195f8e4e04ba2d3968b216e9a6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "548b385c6cf740bbbe0ecd3cd89c5845", "version_major": 2, "version_minor": 0}
```

E1: TrBAcc=0.3257, ValBAcc=0.4351, ValF1=0.1016

```
{"model_id": "9de972fc10694d53839da8e516b06864", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a779f30d32a741fcbff7c94f1ee32e94", "version_major": 2, "version_minor": 0}
```

E2: TrBAcc=0.4241, ValBAcc=0.4624, ValF1=0.1564

```
{"model_id": "27e77172c32e462bbb2dfc36283332c0", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "478a4260ed034e6b967b50dcfd0a2c52", "version_major": 2, "version_minor": 0}
```

E3: TrBAcc=0.4622, ValBAcc=0.5316, ValF1=0.2238

```
{"model_id": "8a560e5424324f0999517252287431f8", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9b5797f0033b44728f1d09a67ffa4584", "version_major": 2, "version_minor": 0}
```

E4: TrBAcc=0.5327, ValBAcc=0.5651, ValF1=0.2891

```
{"model_id": "239e893979804fc3b0b2c03fdb4e6c64", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1797a516237b4d55b6bba25939db7469", "version_major": 2, "version_minor": 0}
```

E5: TrBAcc=0.5132, ValBAcc=0.5467, ValF1=0.2480

```
{"model_id": "3edc88e3dd4f47b6ba7f19182d6fbc98", "version_major": 2, "version_minor": 0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

if w.is_alive():

^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'

^^

AssertionError: can only test a child process

IOStream.flush timed out

```
{"model_id": "42df03bc6e5f4056bf82c5eb6be4dcf3", "version_major": 2, "version_minor": 0}
```

E6: TrBAcc=0.5542, ValBAcc=0.5707, ValF1=0.3145

```
{"model_id": "0a3f56e308fc4491a5f1166c4bd2ae91", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6ddb7522fd554151ad0943f8925ca30c", "version_major": 2, "version_minor": 0}
```

E7: TrBAcc=0.5984, ValBAcc=0.6279, ValF1=0.3772

```
{"model_id": "0e397db78cfd4ab88e1de290ff2f28ac", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c5b4aecb79c446eeb803319b222955a9", "version_major": 2, "version_minor": 0}
```

E8: TrBAcc=0.6256, ValBAcc=0.6292, ValF1=0.3679

```
{"model_id": "c22abdb90cc444278f1e2e788b7211c1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "79a2c14678b64bb1b279b748af6c1632", "version_major": 2, "version_minor": 0}
```

E9: TrBAcc=0.5780, ValBAcc=0.5929, ValF1=0.3145

```
{"model_id": "4046405165e14af3b6c59dcb90fd89c2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e4cc5e048f7d4006a3ebcc958c36f8ef", "version_major": 2, "version_minor": 0}
```

E10: TrBAcc=0.5829, ValBAcc=0.6022, ValF1=0.3465

```
{"model_id": "00fd8fef57674567b2e934ae1dc26d9f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a637a2179cac4ee2adf33056b19ad9cf", "version_major": 2, "version_minor": 0}
```

E11: TrBAcc=0.6095, ValBAcc=0.6273, ValF1=0.3790

```
{"model_id": "c5e5afaa19f44d1c9fc6047efced0194", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c1a2216407bd480f8d00f9de2390fe22", "version_major": 2, "version_minor": 0}
```

E12: TrBAcc=0.6362, ValBAcc=0.6546, ValF1=0.4037

```
{"model_id": "3030afe451d94a2d82cc9666dd3e1c94", "version_major": 2, "version_minor": 0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

```

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "94cfc6d9a07b42fd92809db55df68e8e", "version_major": 2, "vers
ion_minor": 0}

E13: TrBAcc=0.6567, ValBAcc=0.6627, ValF1=0.4361

{"model_id": "88e78937ca484a1f9cbba7d549a6e79b", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "24f96f6c55d8460986ee165df8d9aee1", "version_major": 2, "vers
ion_minor": 0}

E14: TrBAcc=0.6755, ValBAcc=0.6765, ValF1=0.4516

{"model_id": "e2b6bd2f6da5476eb011150d07542cbf", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "d8764339497147fbb08cf00c3551c5e4", "version_major": 2, "vers
ion_minor": 0}

E15: TrBAcc=0.6874, ValBAcc=0.6839, ValF1=0.4867

```

```
{"model_id": "4dde0348220940b5a7a1d4c7a277bb05", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "431ae4916c394677923bce48e5a03c47", "version_major": 2, "version_minor": 0}
```

E16: TrBAcc=0.6897, ValBAcc=0.6789, ValF1=0.4805

```
{"model_id": "247a26b054d44148aeb01453aa2c445d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4aa863752c65403497b0c8898e39ee39", "version_major": 2, "version_minor": 0}
```

E17: TrBAcc=0.6489, ValBAcc=0.6536, ValF1=0.4031

```
{"model_id": "4a546512eaa244688b318d6726cee599", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "46f2dffe3b3a49b186c1a11409b1ceba", "version_major": 2, "version_minor": 0}
```

E18: TrBAcc=0.6391, ValBAcc=0.6549, ValF1=0.4196

```
{"model_id": "eb67a27b491646ab8b088ed13c4ac19c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a41399bb8d9d4a1b8ffcd5abbac269", "version_major": 2, "version_minor": 0}
```

E19: TrBAcc=0.6465, ValBAcc=0.6575, ValF1=0.4075

```
{"model_id": "51f9c471888341d89b09c53286dfd612", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "cfe90c3cfc984ebba0a29949a6c61f0d", "version_major": 2, "version_minor": 0}
```

E20: TrBAcc=0.6674, ValBAcc=0.6816, ValF1=0.4426

```
{"model_id": "39dffc2c06fa452abfefb3d752e57ca7", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "205d7e80d1f84592bcf8ed2a44fa83b4", "version_major": 2, "version_minor": 0}
```

E21: TrBAcc=0.6669, ValBAcc=0.6678, ValF1=0.3998

```
{"model_id": "09f7cb962bc14faf9ad274e653029704", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "24739d3aead5461f8017ad5d9affbf6c", "version_major": 2, "version_minor": 0}
```

E22: TrBAcc=0.6770, ValBAcc=0.6799, ValF1=0.4580

```
{"model_id": "9a468233c9c345c089990c6690740f98", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "be91664a8fc44b30856d67b17529c889", "version_major": 2, "version_minor": 0}
```

E23: TrBAcc=0.6956, ValBAcc=0.6865, ValF1=0.4859

```
{"model_id": "3498116fcc7a4689bda6e791ccda2942", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "eb16d444bf14409c85a8e28822e09222", "version_major": 2, "version_minor": 0}
```

E24: TrBAcc=0.7072, ValBAcc=0.6785, ValF1=0.4894

```
{"model_id": "dbd275dbe720457aa278e1eeb6f0cf16", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5e10e599e90444ec8ca01fe05c777b9e", "version_major": 2, "version_minor": 0}
```

E25: TrBAcc=0.7159, ValBAcc=0.7076, ValF1=0.5167

```
{"model_id": "e288f9763b84433db51a4e97b4412b35", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "eb059df4869d4c0690fa9cb3c03d842b", "version_major": 2, "version_minor": 0}
```

E26: TrBAcc=0.7277, ValBAcc=0.6937, ValF1=0.4992

```
{"model_id": "65f4da63c28e4256bc5116e2a078b6ba", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "18a16333033e4dbfa4b99cea8b93faf8", "version_major": 2, "version_minor": 0}
```

E27: TrBAcc=0.7347, ValBAcc=0.7086, ValF1=0.5199

```
{"model_id": "dcff3c64d25041518bbe7a0a5efd3e02", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "46d9c4fbd10c473bae7a84f1742483ff", "version_major": 2, "version_minor": 0}
```

E28: TrBAcc=0.7471, ValBAcc=0.7026, ValF1=0.5339

```
{"model_id": "03b9cc380493429ebda378b86fd8a190", "version_major": 2, "version_minor": 0}
```



```
{"model_id": "950b0037c9db41bb94d27a406ddcbb36", "version_major": 2, "version_minor": 0}
```

E29: TrBAcc=0.7456, ValBAcc=0.7136, ValF1=0.5532

```
{"model_id": "cb78bbec55a04d10adfc98bc54bb3eab", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "28594b026174410180e3499f21dc0424", "version_major": 2, "version_minor": 0}
```

E30: TrBAcc=0.7609, ValBAcc=0.7105, ValF1=0.5471

```
{"model_id": "32dd3df94e9545638085f4c5851a1c1f", "version_major": 2, "version_minor": 0}
```

TTA Results: Acc=0.5174, BAcc=0.7268, F1=0.5662

```
{"model_id": "1ee31bf37ca64236a55f135599e9ba7a", "version_major": 2, "version_minor": 0}
```

Fold 2 Complete: BAcc=0.7268, F1=0.5662

```
=====
FOLD 3/5
=====
```

```
{"model_id": "8304eea237e64dafbdf0b5b7005003cb", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6f5e890ee35c48269db8c826ca734035", "version_major": 2, "version_minor": 0}
```

E1: TrBAcc=0.3440, ValBAcc=0.4122, ValF1=0.1100

```
{"model_id": "effdec99e2ac4bf4a9d7a4075ca135bc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d87a417a89ee4dfcbeed08e83d5b2142", "version_major": 2, "version_minor": 0}
```

E2: TrBAcc=0.4317, ValBAcc=0.4896, ValF1=0.1705

```
{"model_id": "36e8b2775fdb49c58cd07cd583afd9e8", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9ae376c2622e4b2999357c7bcf10467a", "version_major": 2, "version_minor": 0}
```

E3: TrBAcc=0.4659, ValBAcc=0.5375, ValF1=0.2179

```
{"model_id":"d8e3603a1092437886acb14d08dd7eee","version_major":2,"version_minor":0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
```

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
```

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

```
{"model_id":"761c8802d7794b158d8ff38992e55026","version_major":2,"version_minor":0}
```

E4: TrBAcc=0.5337, ValBAcc=0.5672, ValF1=0.2603

```
{"model_id":"9a8116c71a7d42dbab85eb3ec63228a8","version_major":2,"version_minor":0}
```

```
{"model_id":"ae078ed5a8d641eca7522a2a74413898","version_major":2,"version_minor":0}
```

E5: TrBAcc=0.5081, ValBAcc=0.5710, ValF1=0.2672

```
{"model_id": "639fb68527f04b75a57c4a0aeb518c6e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9207fddd5c6e4b0f9df5c746368f433d", "version_major": 2, "version_minor": 0}
```

E6: TrBAcc=0.5491, ValBAcc=0.6027, ValF1=0.3189

```
{"model_id": "75d12858de294ff68ffd4e9863384970", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "524be5be5ff74d80bf0cb51e6d0ba49d", "version_major": 2, "version_minor": 0}
```

E7: TrBAcc=0.5981, ValBAcc=0.6381, ValF1=0.3729

```
{"model_id": "51fbf8a64ef945c1bf8162dcfa2a9c2b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6ed547fbe7904cbe8d2a38d7570bb92f", "version_major": 2, "version_minor": 0}
```

E8: TrBAcc=0.6245, ValBAcc=0.6342, ValF1=0.3777

```
{"model_id": "e3f19b2a90a3413ea0fdda81656b4165", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a90fdcfcc986452fb6b777ae5f0f99d9", "version_major": 2, "version_minor": 0}
```

E9: TrBAcc=0.5722, ValBAcc=0.5923, ValF1=0.2843

```
{"model_id": "25288428ace841958d003d3ef016d31e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d8eeb1bc482045b395e4d2c5308c56b7", "version_major": 2, "version_minor": 0}
```

E10: TrBAcc=0.5822, ValBAcc=0.5994, ValF1=0.3221

```
{"model_id": "6e80078f8bb142549793ed07f16653ac", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "41db40d124734a83bbaac54abc6f5770", "version_major": 2, "version_minor": 0}
```

E11: TrBAcc=0.6062, ValBAcc=0.6420, ValF1=0.3634

```
{"model_id": "823ebec6f54e4f71b6796da0e3e6714f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c6d7987bd63446ea9328837dce29ad6e", "version_major": 2, "version_minor": 0}
```

E12: TrBAcc=0.6332, ValBAcc=0.6451, ValF1=0.4066

```
{"model_id": "1c94057f618f453aae4fad29b735f4ce", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "afd29fdc0e3948a39c3b37a7e181ef21", "version_major": 2, "version_minor": 0}
```

E13: TrBAcc=0.6549, ValBAcc=0.6502, ValF1=0.4313

```
{ "model_id": "c2b97c7cbd674acfa3696ab9532f9ecc", "version_major": 2, "version_minor": 0 }
```

```
{"model_id": "lafe791c75174941a2d3daf36754f94e", "version_major": 2, "version_minor": 0}
```

E14: TrBAcc=0.6783, ValBAcc=0.6850, ValF1=0.4999

```
{"model_id": "54f3956d761841cda2708139ad3e63ce", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "53037a37403f4ec19bb16bfb1822f28e", "version_major": 2, "version_minor": 0}
```

E15: TrBAcc=0.6969, ValBAcc=0.6935, ValF1=0.5056

```
{"model_id": "c45b46b60bca46d4ac95a48e106b7d2e", "version_major": 2, "version_minor": 0}
```

```
IOStream.flush timed out
```

IOStream.flush timed out

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
```

Traceback (most recent call last):

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self.shutdown_workers()
```

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is alive
```

```
assert self._parent_pid == os.getpid(), 'can only test a child process'
```

```
AssertionError: can only test a child process
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__  
at 0x728d383977e0>
```

```
Traceback (most recent call last):
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1664, in __del__  
    self._shutdown_workers()
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1647, in _shutdown_workers  
    if w.is_alive():  
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in  
is_alive  
    assert self._parent_pid == os.getpid(), 'can only test a child  
process'  
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

```
AssertionError: can only test a child process
```

```
IOStream.flush timed out
```

```
{"model_id": "a34a1fe148c94a35b685531fbcf1f692", "version_major": 2, "vers  
ion_minor": 0}
```

```
E16: TrBAcc=0.6966, ValBAcc=0.6923, ValF1=0.5004
```

```
{"model_id": "509f3473eda741c6a78240a742239e17", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "0ab35f8a86874117a6c351321a326be9", "version_major": 2, "vers  
ion_minor": 0}
```

```
E17: TrBAcc=0.6431, ValBAcc=0.6320, ValF1=0.3371
```

```
{"model_id": "6075589662414b28b95abec57438f554", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "a3a67fcd4b1740dd8d613929959559b0", "version_major": 2, "vers  
ion_minor": 0}
```

```
E18: TrBAcc=0.6293, ValBAcc=0.6389, ValF1=0.3936
```

```
{"model_id": "059dff3104ea4f08902d5799f0f07867", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "ec60d605c604414e990580ee83cf490c", "version_major": 2, "vers  
ion_minor": 0}
```

```
E19: TrBAcc=0.6500, ValBAcc=0.6294, ValF1=0.3647
```

```
{"model_id": "3633717fdcdb42f7be62e3c416accc6d", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "165b6bf552764d3b9d047862c6a73628", "version_major": 2, "version_minor": 0}
```

E20: TrBAcc=0.6639, ValBAcc=0.6707, ValF1=0.4199

```
{"model_id": "b57a98b114c54479b3053d0a98e1b535", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "275455e5bfd043978ab03fc0cf192df2", "version_major": 2, "version_minor": 0}
```

E21: TrBAcc=0.6749, ValBAcc=0.6900, ValF1=0.4336

```
{"model_id": "6fbd4b59e554475f8bb4bd8620fa617e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "bba2b5276e5041b583f924d42053ce3f", "version_major": 2, "version_minor": 0}
```

E22: TrBAcc=0.6727, ValBAcc=0.6734, ValF1=0.4827

```
{"model_id": "cd23ecaf0a274f12b81e5206ad9ed459", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b377050f1e97489ba10fc45310bf4f4d", "version_major": 2, "version_minor": 0}
```

E23: TrBAcc=0.6908, ValBAcc=0.6891, ValF1=0.4651

```
{"model_id": "ec5cbea2d321468c91118e562b64268b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b3c538f30ce845aa9abd92b44d2c6edb", "version_major": 2, "version_minor": 0}
```

E24: TrBAcc=0.7026, ValBAcc=0.7035, ValF1=0.5091

```
{"model_id": "4de569a9bd824250aa39db1f7dba2c3a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7d79930fafb54e04b9aeb171083ecb8", "version_major": 2, "version_minor": 0}
```

E25: TrBAcc=0.7222, ValBAcc=0.7159, ValF1=0.5285

```
{"model_id": "ac7628b850c84e9b8eebf3a43f7e90cf", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "dea23537100c4f83bc1f40b5a9d07453", "version_major": 2, "version_minor": 0}
```

E26: TrBAcc=0.7298, ValBAcc=0.7235, ValF1=0.5360

```
{"model_id": "7d7c437f954849fbb0e7bbcd933e2f47", "version_major": 2, "version_minor": 0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

```
{"model_id": "2f4623fc6fd94f5b936c0d57632712db", "version_major": 2, "version_minor": 0}
```

E27: TrBAcc=0.7333, ValBAcc=0.7213, ValF1=0.5250

```
{"model_id": "595c80e8f1cf42438f9f3faae7dba502", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "30a46efa6f0f43c3b9feaa89ad5217c1", "version_major": 2, "version_minor": 0}
```

E28: TrBAcc=0.7447, ValBAcc=0.7353, ValF1=0.5594

```
{"model_id": "631db2a0ebe949bd8e202efa2e930740", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a323804a20ce41caac8ba4ae4b7a87df", "version_major": 2, "version_minor": 0}
```

E29: TrBAcc=0.7498, ValBAcc=0.7363, ValF1=0.5720

```
{"model_id": "f9f53b4955c2413bb6a72c0b0e414050", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5458a490f602448ab25cb2a9c1f6325c", "version_major": 2, "version_minor": 0}
```

E30: TrBAcc=0.7518, ValBAcc=0.7380, ValF1=0.5851

```
{"model_id": "075aaf1c2a8b4c33a3afca402b6aba5f", "version_major": 2, "version_minor": 0}
```

TTA Results: Acc=0.4907, BAcc=0.7432, F1=0.5891

```
{"model_id": "f744266b40e04884ae790e8f90daeb31", "version_major": 2, "version_minor": 0}
```

Fold 3 Complete: BAcc=0.7432, F1=0.5891

```
=====
FOLD 4/5
=====
```

```
{"model_id": "162be4fa62074cdd9ba8197906ed87d4", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b03e3feb5afa4afeb2a9213b0e64b2e8", "version_major": 2, "version_minor": 0}
```

E1: TrBAcc=0.3356, ValBAcc=0.4378, ValF1=0.1200

```
{"model_id": "a5be0adfc0354de899948a0fa3bf4c24", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "526e1776e1f740c2ac9b26e41c1f7a25", "version_major": 2, "version_minor": 0}
```

E2: TrBAcc=0.4326, ValBAcc=0.4937, ValF1=0.1734

```
{"model_id": "4b417f9f23354592ba32677ebd0606df", "version_major": 2, "version_minor": 0}
```



```
{"model_id": "86de4ae4c5ea455b8276c3173a92d558", "version_major": 2, "version_minor": 0}
```

E3: TrBAcc=0.4677, ValBAcc=0.5478, ValF1=0.2452

```
{"model_id": "e14c853c19ac499fbc22dc05e427c3f4", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3fbf779b6a6944a3846aad70175c1d15", "version_major": 2, "version_minor": 0}
```

E4: TrBAcc=0.5367, ValBAcc=0.5959, ValF1=0.2994

```
{ "model_id": "e69bd0aa16fd4cd6a57dcfdb709ba67f", "version_major": 2, "version_minor": 0 }
```

```
{"model_id": "b9e2a264ecf744de8f97bb1e4572160b", "version_major": 2, "version_minor": 0}
```

E5: TrBAcc=0.5138, ValBAcc=0.5400, ValF1=0.2474

```
{"model_id": "7e3a5896041349c293462fc282564ede", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f38996e3ef5343529bd9b3aae30fd703", "version_major": 2, "version_minor": 0}
```

E6: TrBAcc=0.5580, ValBAcc=0.6132, ValF1=0.3360

```
{"model_id": "3c3c3c0ce0834a2b932725154b80b728", "version_major": 2, "version_minor": 0}
```

```
IOStream.flush timed out
```

```
IOStream.flush timed out
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
```

Traceback (most recent call last):

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self.shutdown_workers()
```

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
```

```
AssertionError: can only test a child process
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__  
at 0x728d383977e0>
```

```
Traceback (most recent call last):
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1664, in __del__  
    self._shutdown_workers()
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1647, in _shutdown_workers  
    if w.is_alive():  
        ^^^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in  
is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child  
process'
```

```
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
AssertionError: can only test a child process
```

```
{"model_id": "367c7249f47547709601ac67f7e5c759", "version_major": 2, "vers  
ion_minor": 0}
```

```
E7: TrBAcc=0.6100, ValBAcc=0.6441, ValF1=0.3827
```

```
{"model_id": "ff2a4d79dc15444ebf1804f8a1a33a22", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "ef6b6c2519994349955bfe42b0d38f10", "version_major": 2, "vers  
ion_minor": 0}
```

```
E8: TrBAcc=0.6279, ValBAcc=0.6486, ValF1=0.3947
```

```
{"model_id": "7f6422a2ed76471fb748039ef285f145", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "8b8de80aabb44ed09813026d91dcb087", "version_major": 2, "vers  
ion_minor": 0}
```

```
E9: TrBAcc=0.5781, ValBAcc=0.6237, ValF1=0.3618
```

```
{"model_id": "0c8d6b3f7c404e54bb6bfeae5a5dbea", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "9ab5f111677149e1b1fed045850b58f4", "version_major": 2, "vers  
ion_minor": 0}
```

```
E10: TrBAcc=0.5863, ValBAcc=0.6699, ValF1=0.3945
```

```
{"model_id": "d11d2ac3437a49e7a6df96ddaa5d7dca", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "382117ba219b4056bd1bfce7d6b42bfa", "version_major": 2, "version_minor": 0}
```

E11: TrBAcc=0.6230, ValBAcc=0.6522, ValF1=0.3783

```
{"model_id": "7924f8ec8c444844b3f88f62ea34b925", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c2d8f8d0c84541b08bac9be94e9fca1a", "version_major": 2, "version_minor": 0}
```

E12: TrBAcc=0.6355, ValBAcc=0.6902, ValF1=0.4526

```
{"model_id": "25cb9c948c5e4c098c80009e0910c211", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "341d7c9b88e14fe395e26281f8ba8fd0", "version_major": 2, "version_minor": 0}
```

E13: TrBAcc=0.6728, ValBAcc=0.6887, ValF1=0.4577

```
{"model_id": "f80c3a0da55141599e395b702bbf3cff", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c7dd0ed023db403c80724ec1d6fdd3cf", "version_major": 2, "version_minor": 0}
```

E14: TrBAcc=0.6808, ValBAcc=0.6806, ValF1=0.4578

```
{"model_id": "40ba8a2a08654b0088f5a075189d878e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "333be8bab07547b9ac43d6cdf014d76a", "version_major": 2, "version_minor": 0}
```

E15: TrBAcc=0.6948, ValBAcc=0.6992, ValF1=0.4847

```
{"model_id": "e41cb609303f4771a492c50e58877db3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f201a2b59ba04402bea872b9f1d4f92b", "version_major": 2, "version_minor": 0}
```

E16: TrBAcc=0.7050, ValBAcc=0.6988, ValF1=0.4797

```
{"model_id": "c0af3b83f59d4cbbb92a5745887a1924", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fb873d3d286c4b7bb0ca3c88e66e54f5", "version_major": 2, "version_minor": 0}
```

E17: TrBAcc=0.6451, ValBAcc=0.6458, ValF1=0.3409

```
{"model_id": "d0bc2927ca414faa9233fb62a38a3baf", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ed2194b9c10c455985fc7e39001013ac", "version_major": 2, "version_minor": 0}
```

E18: TrBAcc=0.6360, ValBAcc=0.6558, ValF1=0.3926

```
{"model_id": "08275282f80a4b4686191f3ead2b2b80", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f84f561f332b4e219cf6a364ec1faddc", "version_major": 2, "version_minor": 0}
```

E19: TrBAcc=0.6526, ValBAcc=0.6829, ValF1=0.3995

```
{"model_id": "67c546ea4d7e4f56bb8a92f4bfdb91c3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "105d407a37554dee8841dddfa31de7aa", "version_major": 2, "version_minor": 0}
```

E20: TrBAcc=0.6642, ValBAcc=0.6750, ValF1=0.4071

```
{"model_id": "52ddedf1a0764f538a6af7f367cd0e02", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c30ddd9af2514405bb3f9e26bf6ff7f0", "version_major": 2, "version_minor": 0}
```

E21: TrBAcc=0.6770, ValBAcc=0.6789, ValF1=0.4277

```
{"model_id": "b36ce1f81e4245a8b49646f58a6f0bf0", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0b8ecbb3106940428396c58efad2f51a", "version_major": 2, "version_minor": 0}
```

E22: TrBAcc=0.6746, ValBAcc=0.6921, ValF1=0.4269

```
{"model_id": "78115fa0c8a04931af8086bcec2b3063", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a04cffe9a05a4b909fa9606a89b39037", "version_major": 2, "version_minor": 0}
```

E23: TrBAcc=0.6877, ValBAcc=0.6970, ValF1=0.4851

```
{"model_id": "007cb4b5a0354502ab5f24d86967f869", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "cd3b341d999946b084075aa35f77424d", "version_major": 2, "version_minor": 0}
```

E24: TrBAcc=0.7033, ValBAcc=0.7076, ValF1=0.5009

```
{"model_id": "02b3977e89e843119778ce3642f37a5b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "822b3ce94ee84ab2a5f3e78c296f8286", "version_major": 2, "version_minor": 0}
```

E25: TrBAcc=0.7187, ValBAcc=0.7252, ValF1=0.4847

```
{"model_id": "5920489699244f04a50dd8b0cfa66b6d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0899f903555f4de0a5fc0b4582772769", "version_major": 2, "version_minor": 0}
```

E26: TrBAcc=0.7327, ValBAcc=0.7186, ValF1=0.5424

```
{"model_id": "e92e6c76bc1243fb9c84ce4b2cbcc58d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a734b4996475426297b3539d267a17bd", "version_major": 2, "version_minor": 0}
```

E27: TrBAcc=0.7418, ValBAcc=0.7324, ValF1=0.5661

```
{"model_id": "197480f9f963410a86845d6a7558d24c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "08555570e5ca4b2d9b7ecb858aabcb00", "version_major": 2, "version_minor": 0}
```

E28: TrBAcc=0.7506, ValBAcc=0.7322, ValF1=0.5838

```
{"model_id": "c0a18229fbfb444cb98eeae4a2242b57", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "feaaaf2108cc3440eae50ab00d4f4409b", "version_major": 2, "version_minor": 0}
```

E29: TrBAcc=0.7579, ValBAcc=0.7295, ValF1=0.5744

```
{"model_id": "1b734fccb6594cc487f4b8f844a6c9a6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6dcee389d22649ed913f4ea6431f8684", "version_major": 2, "version_minor": 0}
```

E30: TrBAcc=0.7593, ValBAcc=0.7392, ValF1=0.5878

```
{"model_id": "68a3b50b38594c978c73bcd932cc3085", "version_major": 2, "version_minor": 0}
```

TTA Results: Acc=0.5221, BAcc=0.7469, F1=0.6025

```
{"model_id": "a7e3c04f76fa425b819abd2e53f154d8", "version_major": 2, "version_minor": 0}
```

Fold 4 Complete: BAcc=0.7469, F1=0.6025

```
=====
FOLD 5/5
=====
```

Warning: You are sending unauthenticated requests to the HF Hub.
Please set a HF_TOKEN to enable higher rate limits and faster
downloads.

```
{"model_id": "d97625f8eed24736b387228cd842f601", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fc10ce65190b4a81afd177c7bd3d2213", "version_major": 2, "version_minor": 0}
```

E1: TrBAcc=0.3283, ValBAcc=0.4002, ValF1=0.0991

```
{"model_id": "70cf9175c90f4ddbafaf368cb3bb5fdf", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "62560a5b735e4df695d7606ba16eb891", "version_major": 2, "version_minor": 0}
```

E2: TrBAcc=0.4230, ValBAcc=0.4923, ValF1=0.1695

```
{"model_id": "5977afadaeae485c9c338b91332bd1ee", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e5168ef92dde456ebb3606069121d52d", "version_major": 2, "version_minor": 0}
```

E3: TrBAcc=0.4640, ValBAcc=0.4952, ValF1=0.2137

```
{"model_id": "acf46fd2e43a4ad6a1cc37ddc6dc253e", "version_major": 2, "version_minor": 0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
```

```

y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
IOStream.flush timed out
IOStream.flush timed out

{"model_id": "f27932543b0149be8ecb456b1a371751", "version_major": 2, "vers
ion_minor": 0}

E4: TrBAcc=0.5297, ValBAcc=0.5749, ValF1=0.2786

{"model_id": "876f73210b864bff9de82c8e22e8c1f6", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "88aa3e468076420a956fb8ca29de0d0f", "version_major": 2, "vers
ion_minor": 0}

E5: TrBAcc=0.5122, ValBAcc=0.5595, ValF1=0.2673

{"model_id": "93ad592318724fc2af90287dd1068e08", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "2270c854f22d4906afaa8f8d0f41941e", "version_major": 2, "vers
ion_minor": 0}

E6: TrBAcc=0.5607, ValBAcc=0.6127, ValF1=0.3326

```

```
{"model_id": "1c7abf97e9c74b5d998610e44a897417", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a8716f2547e8490189ff6b1eca438335", "version_major": 2, "version_minor": 0}
```

E7: TrBAcc=0.5963, ValBAcc=0.6371, ValF1=0.3696

```
{"model_id": "6edeafecb9d841a8980b6922fe000743", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d8c96bef0a9b4fff91ff4518c6e6ce2e", "version_major": 2, "version_minor": 0}
```

E8: TrBAcc=0.6177, ValBAcc=0.6335, ValF1=0.3738

```
{"model_id": "df0474c7efcc484095228adc47905d57", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "893d767b9c9a478ebd2a3fb5cc55fc17", "version_major": 2, "version_minor": 0}
```

E9: TrBAcc=0.5803, ValBAcc=0.6010, ValF1=0.3456

```
{"model_id": "2225fe72835845e4bc3fc4e7f0327f03", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f6e037b8803c4e9c93678998a511d3e0", "version_major": 2, "version_minor": 0}
```

E10: TrBAcc=0.5841, ValBAcc=0.5993, ValF1=0.3276

```
{"model_id": "a3dbde75cacc4b329e019097f0dddb1c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c0dbeffeb6874df28287b0a1aeed8747", "version_major": 2, "version_minor": 0}
```

E11: TrBAcc=0.6161, ValBAcc=0.6246, ValF1=0.3694

```
{"model_id": "ed9d61bc18004c5cbb3238e636962c16", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "22f6c43ff4da4f0aa09d25da92c20aa5", "version_major": 2, "version_minor": 0}
```

E12: TrBAcc=0.6339, ValBAcc=0.6324, ValF1=0.3616

```
{"model_id": "f940fea64df94010b3bc2b9ca236bf9a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2b31de4958274ed091d438bda46fe1ba", "version_major": 2, "version_minor": 0}
```


E13: TrBAcc=0.6603, ValBAcc=0.6658, ValF1=0.4269

```
{"model_id": "e5df78268ef2427c8b0636cec26e7ae3", "version_major": 2, "version_minor": 0}
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
```

Traceback (most recent call last):

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
self._shutdown_workers()
```

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown workers
```

```
if w.is_alive():
```

^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is alive
```

```
assert self._parent_pid == os.getpid(), 'can only test a child process'
```

[illegible]

```
AssertionError: can only test a child process
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
```

Traceback (most recent call last):

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in del
```

```
self.shutdown_workers()
```

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown workers
```

```
if w.is alive():
```

—

^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is alive
```

```
assert self._parent_pid == os.getpid(), 'can only test a child process'
```

^^^

```
AssertionError: can only test a child process
```

```
{"model_id": "2d18a4a99f3b4d258017e175c0e696d3", "version_major": 2, "version_minor": 0}
```

E14: TrBAcc=0.6742, ValBAcc=0.6766, ValF1=0.4420

```
{"model_id": "a8e00c15e1d14be092885c33593f73f6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "85f6d132b89d44d084a85e2f0d7e5db4", "version_major": 2, "version_minor": 0}
```

E15: TrBAcc=0.6921, ValBAcc=0.6911, ValF1=0.4772

```
{"model_id": "b99c82d70f7b4fa6b93d84fbb48a6cad", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "bbbc7000e0c24045ae9a5a3b11f74a04", "version_major": 2, "version_minor": 0}
```

E16: TrBAcc=0.6987, ValBAcc=0.6891, ValF1=0.4820

```
{"model_id": "1ae7266f13bf422db7015dcf7a25af68", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "34f79ea831834b958115ceeae9fbf1f", "version_major": 2, "version_minor": 0}
```

E17: TrBAcc=0.6502, ValBAcc=0.6642, ValF1=0.4061

```
{"model_id": "a1343412872443eb8b4b2212dd786d2c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3c29e20d9c6d40068ff953e86fc63476", "version_major": 2, "version_minor": 0}
```

E18: TrBAcc=0.6350, ValBAcc=0.6468, ValF1=0.3961

```
{"model_id": "cff54a05e0804d5cad1aee139be80d5f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "88458ce04d504f238993f7ada03ca4c8", "version_major": 2, "version_minor": 0}
```

E19: TrBAcc=0.6567, ValBAcc=0.6799, ValF1=0.4245

```
{"model_id": "da5b9b5186994e64beebda3c5572042e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6274a09fa9534024bbcb4e3c1dd24af5", "version_major": 2, "version_minor": 0}
```

E20: TrBAcc=0.6723, ValBAcc=0.6475, ValF1=0.4106

```
{"model_id": "dd87d4666b9d41a19945dd261f505ea4", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a140b1b16c4f43e8aae50bde63176fc1", "version_major": 2, "version_minor": 0}
```

E21: TrBAcc=0.6733, ValBAcc=0.6529, ValF1=0.4128

```
{"model_id": "6ab1f5e772654d1cb5f9300cb3dd2bad", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b891dde709c6463d9a35cde67a8b15eb", "version_major": 2, "version_minor": 0}
```

E22: TrBAcc=0.6838, ValBAcc=0.6930, ValF1=0.4403

```
{"model_id": "fef8811a67e54684b52a0900cf45ce3c", "version_major": 2, "version_minor": 0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

if w.is_alive():

^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'

^^

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

if w.is_alive():

^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'

^^

AssertionError: can only test a child process

IOStream.flush timed out
IOStream.flush timed out

```
{"model_id":"15e1c454f80e49129d5cdfed034d2bc9","version_major":2,"version_minor":0}
```

E23: TrBAcc=0.6930, ValBAcc=0.7046, ValF1=0.4926

```
{"model_id":"c63b279eaaaa469ca84485a30dd84fc8","version_major":2,"version_minor":0}
```

```
{"model_id":"1c34bf2ff5284114a6d9aa7f9f4fb490","version_major":2,"version_minor":0}
```

E24: TrBAcc=0.7052, ValBAcc=0.6893, ValF1=0.5060

```
{"model_id":"98e2deec9eb54ff1a9404d968a3d9720","version_major":2,"version_minor":0}
```

```
{"model_id":"5399bd3d208841cda600cf2dd97acff8","version_major":2,"version_minor":0}
```

E25: TrBAcc=0.7171, ValBAcc=0.6906, ValF1=0.4972

```
{"model_id":"5c16a04f20254fec8ea489f91b1908d7","version_major":2,"version_minor":0}
```

E26: TrBAcc=0.7319, ValBAcc=0.7049, ValF1=0.5149

```
{"model_id":"ce4c0dcab99e437a829e8aecc7111875","version_major":2,"version_minor":0}
```

```
{"model_id":"e7269d609577416bb5be88056aac0bd2","version_major":2,"version_minor":0}
```

E27: TrBAcc=0.7359, ValBAcc=0.6990, ValF1=0.5146

```
{"model_id":"cb7d908b4ea74aabb6385be5aef82415","version_major":2,"version_minor":0}
```

```
{"model_id":"4ae4d7ce0219459aa2045d0a966b3b10","version_major":2,"version_minor":0}
```

E28: TrBAcc=0.7483, ValBAcc=0.7168, ValF1=0.5391

```
{"model_id":"5a84adf03de249e49c0599d161bc0eb6","version_major":2,"version_minor":0}
```

```
{"model_id":"09d8c009124d4c389cf2ee12b5a11cdb","version_major":2,"version_minor":0}
```

E29: TrBAcc=0.7548, ValBAcc=0.7159, ValF1=0.5752

```
{"model_id": "95bac629e80d4eea9049d9381b4a7022", "version_major": 2, "version_minor": 0}
```

IOStream.flush timed out

IOStream.flush timed out

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
```

Traceback (most recent call last):

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
self._shutdown_workers()
```

File

```
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown workers
```

```
if w.is_alive():
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
```

```
assert self._parent_pid == os.getpid(), 'can only test a child
cess'
```

^^^

```
AssertionError: can only test a child process
```

```
{"model_id": "64229431c95d4a83bc11efbebac5cdac", "version_major": 2, "version_minor": 0}
```

E30: TrBAcc=0.7580, ValBAcc=0.7244, ValF1=0.5661

```
{"model_id": "ee3bef04963a41c98c3553b0f18d6859", "version_major": 2, "version_minor": 0}
```

TTA Results: Acc=0.5156, BAcc=0.7334, F1=0.5783

```
{"model_id": "857890e8aa80488794bdc500a0f9f3ca", "version_major": 2, "version_minor": 0}
```

Fold 5 Complete: BAcc=0.7334, F1=0.5783

5-FOLD CV RESULTS

Accuracy: 0.5051 ± 0.0168

Balanced Accuracy: 0.7335 ± 0.0108

Macro F1: 0.5817 ± 0.0129

95% CI - Accuracy: [0.4818, 0.5284]

95% CI - Balanced Acc: [0.7186, 0.7485]

95% CI - Macro F1: [0.5638, 0.5995]

```

# CELL 20B: CLINICAL SAFETY (MEL operating point + uncertainty
# referral + calibration)

import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix, roc_curve,
roc_auc_score, precision_recall_fscore_support

MEL_CLASS = CONFIG["class_names"].index("MEL")

def mel_operating_point(y_true, y_probs, target_sens=0.95):
    """
    Choose a MEL probability threshold to reach >= target sensitivity
    (recall for MEL),
    and report specificity, precision, F1 at that threshold.
    """
    mel_true = (y_true == MEL_CLASS).astype(int)
    mel_score = y_probs[:, MEL_CLASS]

    fpr, tpr, thr = roc_curve(mel_true, mel_score)
    # tpr = sensitivity, fpr = 1 - specificity
    idx = np.where(tpr >= target_sens)[0]
    if len(idx) == 0:
        best_idx = np.argmax(tpr)
    else:
        best_idx = idx[np.argmin(fpr[idx])] # highest specificity
        among those meeting target sensitivity

    threshold = thr[best_idx]
    mel_pred = (mel_score >= threshold).astype(int)

    tn, fp, fn, tp = confusion_matrix(mel_true, mel_pred).ravel()
    sens = tp / (tp + fn + 1e-9)
    spec = tn / (tn + fp + 1e-9)
    prec = tp / (tp + fp + 1e-9)
    f1 = 2 * prec * sens / (prec + sens + 1e-9)
    auc = roc_auc_score(mel_true, mel_score) if 0 < mel_true.sum() <
len(mel_true) else 0.0

    return {
        "target_sensitivity": target_sens,
        "chosen_threshold": float(threshold),
        "MEL_sensitivity": float(sens),
        "MEL_specificity": float(spec),
        "MEL_precision": float(prec),
        "MEL_F1": float(f1),
        "MEL_AUROC": float(auc),
        "tp": int(tp), "fp": int(fp), "tn": int(tn), "fn": int(fn),
    }

```

```

def referral_policy(y_true, y_pred, y_probs, uncertainties,
reject_percentiles=(80, 90, 95)):
    """
    Uncertainty-based referral: reject top X% most uncertain cases,
    compute coverage + performance on remaining cases.
    """
    results = []
    uncertainties = np.asarray(uncertainties)
    for p in reject_percentiles:
        thr = np.percentile(uncertainties, p)
        keep = uncertainties <= thr
        coverage = keep.mean()

        yt = y_true[keep]
        yp = y_pred[keep]
        ypr = y_probs[keep]

        acc = (yp == yt).mean() if len(yt) else 0
        # MEL sensitivity on kept cases
        mel_true = (yt == MEL_CLASS).astype(int)
        mel_pred = (yp == MEL_CLASS).astype(int)
        tn, fp, fn, tp = confusion_matrix(mel_true, mel_pred,
labels=[0,1]).ravel()
        mel_sens = tp / (tp + fn + 1e-9)

        results.append({
            "reject_percentile": p,
            "uncertainty_threshold": float(thr),
            "coverage": float(coverage),
            "accuracy_on_kept": float(acc),
            "MEL_sensitivity_on_kept": float(mel_sens),
            "n_kept": int(len(yt)),
            "n_referred": int((~keep).sum())
        })
    return pd.DataFrame(results)

# Aggregate across folds using stored fold_results
# NOTE: fold_results[i]['labels'], ['preds'], ['probs'] exist in your
notebook.
all_y_true = np.concatenate([r["labels"] for r in fold_results])
all_y_pred = np.concatenate([r["preds"] for r in fold_results])
all_y_prob = np.concatenate([r["probs"] for r in fold_results])

# If you computed uncertainties per fold (from estimate_uncertainty),
collect them here.
# If not available, you can re-run uncertainty on the val loaders per
fold or skip referral section.
# For now, we'll recompute uncertainties on a quick subset by re-using
the first fold model + val data if you saved them.
# (If you want full referral analysis per-fold, tell me and I'll give

```

```

the exact fold-level code.)
print("Clinical safety: MEL operating point on pooled CV predictions")
mel_report = mel_operating_point(all_y_true, all_y_prob,
target_sens=0.95)
mel_df = pd.DataFrame([mel_report])
print(mel_df.to_string(index=False))
wandb.log({"clinical/MEL_operating_point":
wandb.Table(dataframe=mel_df)})

```

```

print("\nTip: For uncertainty-based referral, you need uncertainties
aligned with these pooled predictions.")
print("If you want, I can provide the exact fold-wise code to store
uncertainties during validation and aggregate them.")

```

```

Clinical safety: MEL operating point on pooled CV predictions
  target_sensitivity  chosen_threshold  MEL_sensitivity
MEL_specificity  MEL_precision  MEL_F1  MEL_AUROC  tp    fp    tn    fn
0.348359          0.95          0.140667          0.950022
0.240591 0.383949  0.819794 4296 13560 7249 226

```

```

-----
-----
Error                                Traceback (most recent call
last)
Cell In[40], line 94
      92 mel_df = pd.DataFrame([mel_report])
      93 print(mel_df.to_string(index=False))
--> 94 wandb.log({"clinical/MEL_operating_point":
wandb.Table(dataframe=mel_df)})
      96 print("\nTip: For uncertainty-based referral, you need
uncertainties aligned with these pooled predictions.")
      97 print("If you want, I can provide the exact fold-wise code to
store uncertainties during validation and aggregate them.")

```

```

File
/usr/local/lib/python3.12/dist-packages/wandb/sdk/lib/preinit.py:36,
in PreInitCallable.<locals>.preinit_wrapper(*args, **kwargs)
      35 def preinit_wrapper(*args: Any, **kwargs: Any) -> Any:
--> 36     raise wandb.Error(f"You must call wandb.init() before
{name}()")

```

```

Error: You must call wandb.init() before wandb.log()

```

```

wandb.log({"comparison_protocol": "Budgeted baselines/ablations: 10
epochs; UHViT full CV: 30 epochs"})

```

```

# CELL 21: ABLATION STUDY (FIXED LABELS + BEST-EPOCH SELECTION)

```

```

print("\n" + "="*50 + "\nABLATION STUDY\n" + "="*50)

```

```

# Define ablation models

```



```

ablation_models = {
    "UHViT (Full)": lambda: UHViT(CONFIG["num_classes"],
CONFIG["dropout_rate"]),
    "Swin-T Only": lambda: SwinOnly(CONFIG["num_classes"],
CONFIG["dropout_rate"]),
    "EfficientNet Only": lambda:
EfficientNetOnly(CONFIG["num_classes"], CONFIG["dropout_rate"]),
    "Concat Fusion": lambda: UHViTConcatFusion(CONFIG["num_classes"],
CONFIG["dropout_rate"]),
}

# Use a fixed fold split for budgeted ablation (fold 1)
train_idx, val_idx = list(skf.split(train_df, all_labels))[0]
abl_train_df, abl_val_df =
train_df.iloc[train_idx].reset_index(drop=True),
train_df.iloc[val_idx].reset_index(drop=True)

# Datasets & loaders (same transforms as main training)
abl_train_dataset = ISIC2019Dataset(abl_train_df, train_img_dir,
get_train_transforms(CONFIG["img_size"]))
abl_val_dataset = ISIC2019Dataset(abl_val_df, train_img_dir,
get_val_transforms(CONFIG["img_size"]))

# FIXED: build exactly one label per sample
abl_train_labels = []
for i in range(len(abl_train_df)):
    row = abl_train_df.iloc[i]
    label = 0
    for j, cls in enumerate(CONFIG["class_names"]):
        if cls in abl_train_df.columns and float(row[cls]) == 1.0:
            label = j
            break
    abl_train_labels.append(label)
abl_train_labels = np.array(abl_train_labels)

abl_sampler = create_balanced_sampler(abl_train_labels)

abl_train_loader = DataLoader(
    abl_train_dataset,
    batch_size=CONFIG["batch_size"],
    sampler=abl_sampler,
    num_workers=CONFIG["num_workers"],
    pin_memory=True,
)

abl_val_loader = DataLoader(
    abl_val_dataset,
    batch_size=CONFIG["batch_size"],
    shuffle=False,
    num_workers=CONFIG["num_workers"],

```

```

        pin_memory=True,
    )

    ablation_results = []

    # Train each ablation model for 10 epochs (budgeted), report BEST val BAcc
    for name, model_fn in ablation_models.items():
        print(f"\nTraining {name}...")
        model = model_fn().to(CONFIG["device"])

        samples_per_class = np.maximum(np.bincount(abl_train_labels,
minlength=CONFIG["num_classes"]), 1)
        criterion = ClassBalancedLoss(samples_per_class,
CONFIG["num_classes"])

        optimizer = AdamW(model.parameters(), lr=CONFIG["lr"],
weight_decay=CONFIG["weight_decay"])
        scheduler = CosineAnnealingWarmRestarts(optimizer, T_0=10,
T_mult=2)

        best_bacc, best_acc, best_f1 = -1.0, 0.0, 0.0

        for epoch in range(10):
            train_epoch(model, abl_train_loader, criterion, optimizer,
scheduler, CONFIG["device"], epoch + 1)
            _, val_acc, val_bacc, val_f1, _, _, _ = validate(model,
abl_val_loader, criterion, CONFIG["device"])

            if val_bacc > best_bacc:
                best_bacc, best_acc, best_f1 = val_bacc, val_acc, val_f1

        print(f"{name} (BEST): Acc={best_acc:.4f}, BAcc={best_bacc:.4f},
F1={best_f1:.4f}")

        ablation_results.append({
            "Model": name,
            "Accuracy_best": float(best_acc),
            "BalancedAcc_best": float(best_bacc),
            "MacroF1_best": float(best_f1),
            "epochs": 10,
            "protocol": "budgeted_single_fold_best_epoch"
        })

    del model
    torch.cuda.empty_cache()

abl_df = pd.DataFrame(ablation_results)
print("\nAblation Results (budgeted, single fold, best epoch):")
print(abl_df.to_string(index=False))

```

```
try:
    wandb.log({"ablation_study": wandb.Table(dataframe=abl_df)})
except Exception as e:
    print("W&B log skipped:", e)
```

```
=====
ABLATION STUDY
=====
```

Training UHViT (Full)...

```
{"model_id": "bdb8bcd365574f6aaed388b208a2f36c", "version_major": 2, "version_minor": 0}

{"model_id": "e5c753cc4bc645d59e2bede5b93f65d6", "version_major": 2, "version_minor": 0}

{"model_id": "de88cc46e38445279b8a64f3bdd3b139", "version_major": 2, "version_minor": 0}

{"model_id": "982829d0f5a64ae293a84fa6f7f73603", "version_major": 2, "version_minor": 0}

{"model_id": "b7990d7a69d442e2bce2bf9b85e61602", "version_major": 2, "version_minor": 0}

{"model_id": "0329af79dbf54e039320b8463c9ea529", "version_major": 2, "version_minor": 0}

{"model_id": "1bb16dc903ed4ba4bc0fdec17a645c70", "version_major": 2, "version_minor": 0}

{"model_id": "cb85ac9edbc149308c04fc9873d30314", "version_major": 2, "version_minor": 0}

{"model_id": "0ece33d6b0384b5bac8c09cdf43ee91f", "version_major": 2, "version_minor": 0}

{"model_id": "d91650ae10a446f2a375c81f952ea072", "version_major": 2, "version_minor": 0}

{"model_id": "157d79b19c7449688ae231eba193769b", "version_major": 2, "version_minor": 0}

{"model_id": "bda9aed004b8426dbf1d24c682a06831", "version_major": 2, "version_minor": 0}

{"model_id": "bd9da9e4b1ce492bab263e28fc986749", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c556467b93b04a13bca0c783114ae04b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fd7806f562c546269a1e0ff3215c2626", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0a7375b022ba49a79df1bd628db30a05", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "92a9ff6c605d4b2bbc569faa2623f163", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2558ec5d25ac4d309c4e64fcf1afa752", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "84dc0608dafd4197b8584c6aa4b929a1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "8b4c1593fb1f4a21ad5aaa6e9a4c37d1", "version_major": 2, "version_minor": 0}
```

UHViT (Full) (BEST): Acc=0.3306, BAcc=0.6317, F1=0.3917

Training Swin-T Only...

```
{"model_id": "6210c6048dd94c069008770858863d3f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a66b005f9f844aaba62ba4913e20c76f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "752814178bcf4b6487fe80eb16885aac", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "30c31d85b71e48e7a31501a105e5d022", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "681f328e4f5b4374b7c5e72f9eb08058", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "22690b9dfc9a4e8594c9cc19b8ab581c", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "c1d41186f3474afc90e718977689a3a2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "99d22a143f89471295c34ac9662bd7f5", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9057e8bf0a0245e2a0ff96215fa48f3f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "89968b5122fb4a4a903a5538d420cdfa", "version_major": 2, "version_minor": 0}
```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__  
at 0x728d383977e0>
```

```
Traceback (most recent call last):
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1664, in __del__  
    self._shutdown_workers()
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1647, in _shutdown_workers  
    if w.is_alive():  
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in  
is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child  
process'
```

```
^^^^^^^^^Exception ignored in: ^<function  
_MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>^^
```

```
^Traceback (most recent call last):
```

```
^ File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1664, in __del__  
^^    ^^self._shutdown_workers()^
```

```
^ File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1647, in _shutdown_workers  
^^    ^if w.is_alive():^  
^ ^ ^ ^ ^ ^ ^ ^
```

```
AssertionError^: ^can only test a child process^
```

```
^^^Exception ignored in: <function  
_MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>^
```

```
^Traceback (most recent call last):
```

```
^ File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1664, in __del__  
^^    ^self._shutdown_workers()
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in  
is_alive
```

```
File  
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p  
y", line 1647, in _shutdown_workers  
    assert self._parent_pid == os.getpid(), 'can only test a child  
process'    if w.is_alive():
```

```
        ^ ^ ^ ^^^^^^^^^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
```

```
is_alive
^      ^assert self._parent_pid == os.getpid(), 'can only test a child
process'^
^ ^      ^ ^      ^ ^      ^ ^      ^ ^      ^ ^      ^ ^      ^ ^      ^ ^
^AssertionError^^: ^can only test a child process^^
^^^^Exception ignored in: ^^<function
_MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>^
^Traceback (most recent call last):
^  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
^^      ^self._shutdown_workers()^

  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
AssertionError:      can only test a child processif w.is_alive():

      ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
      ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "d6ca1c59f3c04ee09b7ba709725258d7", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "f29f795ea6e741129eac8254bd88977a", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "f25cb1460ed04ffb907034ec4510e06b", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "3f0ff288852d415b860d04b7dda46d0e", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "ed5d4eb7e2e947708394d90db3a77e0f", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "d642002ea976423ebbb4a2fb02ba545d", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "6b7cf97e258d49c0938a0ac886d51392", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "5bcf71c692124b329cf9ac6079f8ed8c", "version_major": 2, "vers
ion_minor": 0}
```

```
{"model_id": "35bea2c47d494fc6aae7274bc2f5cb47", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "354d5100791f4ad79b1f9ba8beadd6d", "version_major": 2, "version_minor": 0}
```

Swin-T Only (BEST): Acc=0.4227, BAcc=0.6745, F1=0.4213

Training EfficientNet Only...

```
{"model_id": "c63290b71da0484593271f5e57f4aff0", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f516e42af7cb4eaa918d1296a04d4169", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "76f129fab57547109a1617ce89fd8192", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9b30f7d4d5694ee58b7acd7ab44db6dc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e8d74400d7be4a4f83c8d954a7256300", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6767506f411c401db0fecc84364ae330", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3ddd6b0d3b8446cabb9948add6e828da", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "3e5496e46f0f4f84a72da6c358ea3409", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1c22de1b3435407a86f4a9d507b6d7a9", "version_major": 2, "version_minor": 0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
    ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child
```

```
process'
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "a4969c7ad1ef49178323d17f6adda3cf", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "67b0adf31a2f4e63978d62f820868377", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "04ebaaa76e1544b4b05253136d9ca3da", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "520bb5da19044bb682bfedacd50d5774", "version_major": 2, "vers
ion_minor": 0}

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
```


[illegible]

```
{"model_id": "50e23337b3394f3ea5371a78c50d6742", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b77bdb4325f146b0843af76bb3636544", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5d3561dddb8e4c879fc058fd5c74858d", "version_major": 2, "version_minor": 0}
```

```
{ "model_id": "5e367f8df1f346649c8a874a10c1448b", "version_major": 2, "version_minor": 0 }
```

```
{"model_id": "143ca2f67c7645f285732e00e96fbb62", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9827ea060c554a2488ad7042a32a1dbc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9760feb0d34244e892fe994635e26532", "version_major": 2, "version_minor": 0}
```

EfficientNet Only (BEST): Acc=0.3570, BAcc=0.6232, F1=0.3358

Training Concat Fusion...

```
{"model_id": "82bd8707a99643a38ae37e58de13b35a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "bc256d2c411d451298fdd13ddbe1b41e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "51bcabb1648140c58cd5060b8bc9a72d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2c65ce058c7445d49a418d32cbe75d70", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "5d4cd1455e5d4387a66ffee5b2de6e92", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "27895fda062847409ee1498cacdd55ec", "version_major": 2, "version_minor": 0}
```

```
{ "model_id": "24d92639cfbd4a48a5ed421e5bb872fd", "version_major": 2, "version_minor": 0 }
```

```
{ "model_id": "02aa1fbf33ed4a54ba013a60f03234b4", "version_major": 2, "version_minor": 0 }
```

```
{"model_id":"cfbe494b6ee04f70901630717fd22b60","version_major":2,"version_minor":0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

if w.is_alive():

^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'

^^

AssertionError: can only test a child process

IOStream.flush timed out

IOStream.flush timed out

```
{"model_id":"107826bc837c41769aa177f4ea576922","version_major":2,"version_minor":0}
```

```
{"model_id":"0f5765c4c8cb425aba604f74da251cc9","version_major":2,"version_minor":0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers

if w.is_alive():

^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'

```

AssertionError: can only test a child process
IOStream.flush timed out
IOStream.flush timed out

{"model_id": "28561d24f7634cc7933fa6351b66b413", "version_major": 2, "version_minor": 0}

{"model_id": "d4256a7767d740c4b1dedb222779e2b9", "version_major": 2, "version_minor": 0}

{"model_id": "027681bfe1aa4442959dc136e1516d26", "version_major": 2, "version_minor": 0}

{"model_id": "ef2013775380454697c7b94a6eab1e1a", "version_major": 2, "version_minor": 0}

{"model_id": "4b19c130b0a14a4ba615d799b253bbc7", "version_major": 2, "version_minor": 0}

{"model_id": "e14c3b8484db4a92aa86157e64f48821", "version_major": 2, "version_minor": 0}

{"model_id": "19268ae743854181b4db30aade11e86c", "version_major": 2, "version_minor": 0}

{"model_id": "e9115fc5834a42148d74119d6e09ced0", "version_major": 2, "version_minor": 0}

{"model_id": "332a8336dbf242209d18a80aaf0e6acb", "version_major": 2, "version_minor": 0}

Concat Fusion (BEST): Acc=0.4579, BAcc=0.6884, F1=0.4658

Ablation Results (budgeted, single fold, best epoch):

```

epochs	Model	Accuracy_best	BalancedAcc_best	MacroF1_best
10	UHViT (Full)	0.330570	0.631712	0.391707
10	budgeted_single_fold_best_epoch			
10	Swin-T Only	0.422735	0.674525	0.421324
10	budgeted_single_fold_best_epoch			
10	EfficientNet Only	0.357016	0.623233	0.335812
10	budgeted_single_fold_best_epoch			
10	Concat Fusion	0.457865	0.688450	0.465830
10	budgeted_single_fold_best_epoch			

```

# CELL 22: SOTA COMPARISON (FIXED: proper samples_per_class + BEST epoch)

print("\n" + "="*50 + "\nSOTA COMPARISON (budgeted single-fold)\n" +
      "="*50)

```

```

# This cell assumes you already ran the FIXED ablation cell (CELL 21)
# that created:
#   abl_train_loader, abl_val_loader, abl_train_labels
# If not, run CELL 21 first.

assert "abl_train_loader" in globals() and "abl_val_loader" in
globals() and "abl_train_labels" in globals(), \
    "Run the fixed Ablation Study cell (CELL 21) first."

# Define criterion inputs correctly for this fold
samples_per_class = np.maximum(np.bincount(abl_train_labels,
minlength=CONFIG["num_classes"]), 1)
criterion = ClassBalancedLoss(samples_per_class,
CONFIG["num_classes"])

sota_results = []

for sota_name in SOTA_MODELS:
    print(f"\nTraining SOTA baseline: {sota_name} ...")
    try:
        model = create_sota_model(sota_name,
CONFIG["num_classes"]).to(CONFIG["device"])
        optimizer = AdamW(model.parameters(), lr=CONFIG["lr"],
weight_decay=CONFIG["weight_decay"])
        scheduler = CosineAnnealingWarmRestarts(optimizer, T_0=10,
T_mult=2)

        best_bacc, best_acc, best_f1 = -1.0, 0.0, 0.0

        for epoch in range(10):
            train_epoch(model, abl_train_loader, criterion, optimizer,
scheduler, CONFIG["device"], epoch + 1)
            _, val_acc, val_bacc, val_f1, _, _, _ = validate(model,
abl_val_loader, criterion, CONFIG["device"])

            if val_bacc > best_bacc:
                best_bacc, best_acc, best_f1 = val_bacc, val_acc,
val_f1

            print(f"{sota_name} (BEST): Acc={best_acc:.4f},
BAcc={best_bacc:.4f}, F1={best_f1:.4f}")

        sota_results.append({
            "Model": sota_name,
            "Accuracy_best": float(best_acc),
            "BalancedAcc_best": float(best_bacc),
            "MacroF1_best": float(best_f1),
            "epochs": 10,
            "protocol": "budgeted_single_fold_best_epoch"
        })

```

```

        del model
        torch.cuda.empty_cache()

    except Exception as e:
        print(f"Error with {sota_name}: {e}")

# Add your model result (from 5-fold CV means) as "Ours"
# These variables accs/baccs/fls exist in your CV summary cell; if
not, you can skip or hardcode.
if "accs" in globals() and "baccs" in globals() and "fls" in
globals():
    sota_results.append({
        "Model": "UHViT (Ours, 5-fold CV mean)",
        "Accuracy_best": float(np.mean(accs)),
        "BalancedAcc_best": float(np.mean(baccs)),
        "MacroF1_best": float(np.mean(fl_s)),
        "epochs": CONFIG.get("epochs", 30),
        "protocol": "5-fold CV_mean"
    })

sota_df = pd.DataFrame(sota_results).sort_values("BalancedAcc_best",
ascending=False)
print("\nSOTA Comparison Results:")
print(sota_df.to_string(index=False))

try:
    wandb.log({"sota_comparison": wandb.Table(dataframe=sota_df)})
except Exception as e:
    print("W&B log skipped:", e)

```

```

=====
SOTA COMPARISON (budgeted single-fold)
=====

```

Training SOTA baseline: resnet50 ...

```

{"model_id": "711d7106498848a4992a3745e4b52a6c", "version_major": 2, "version_minor": 0}

{"model_id": "b39b6ff57fb44626893f362d7ff26d96", "version_major": 2, "version_minor": 0}

{"model_id": "3174b0812f854bbaaa6f942cca2b44db", "version_major": 2, "version_minor": 0}

{"model_id": "04c54cd648694523b134f032594eba59", "version_major": 2, "version_minor": 0}

```

```
{"model_id": "4ea1bce4717f41e682f82b08c22b2f88", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "837bb7dbc89843d4b208a37e7da4e3dc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d6c203bb5cfd4f6fb58cc1fba095a588", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "aa23ec0f009c4b03bb64d8edb4b96834", "version_major": 2, "version_minor": 0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
```

```
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
```

```
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

```

Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ~~~~~~
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ~~~~~~
AssertionError: can only test a child process

{"model_id": "ale5f1d3e2b9419894f0e77b93ec9b64", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "9c4e2c2ef9ea4b59b24d0f3634418380", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "e5b06ab5599d424e956f61d565bb5c46", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "ec98674f40594a2f86f138796a6329a2", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "3e9db4bc7bb74be883fadb2f0c1cff5b", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "62c7e7a454ed498f9ecd18ede275e170", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "ed87876dd52e4d0eb4c46e413c3b9a13", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "28f9b40dca124d8d851a3ac946214afb", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "c181882d097947aa8d34dac993212ca8", "version_major": 2, "vers
ion_minor": 0}

IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):

```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
Exception ignored in:      <function
_MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>
w.is_alive():
Traceback (most recent call last):
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
    ^ ^ ^ ^ ^ ^ ^ ^ ^^^^^^^^^^^
^ File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
^   ^assert self._parent_pid == os.getpid(), 'can only test a child
process'^
^ ^ ^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^ ^ ^^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^
```



```

^AssertionError^: ^can only test a child process^
^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out

{"model_id": "d23df4df2c0a46c1a9539494bc1265e7", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "5a241c6a8f2a44dba21595db9e67d341", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "f6e7da158387493fab6a240d8a856465", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "1dd4278479a9450782019d3c120a109f", "version_major": 2, "vers
ion_minor": 0}

resnet50 (BEST): Acc=0.1508, BAcc=0.4836, F1=0.1594

Training SOTA baseline: efficientnet_b3 ...

{"model_id": "eaf24a8f7d3044298667d00799761086", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "fa37a84159194781b621fccb22bdd8bc", "version_major": 2, "vers
ion_minor": 0}

```

```
{"model_id": "19a2b00676a4495ba1cb08dcf1ed56f2", "version_major": 2, "version_minor": 0}

{"model_id": "b310270940934b85b6f3129bb017f9e7", "version_major": 2, "version_minor": 0}

{"model_id": "c778bca6e081475488168a07c9df277c", "version_major": 2, "version_minor": 0}

{"model_id": "ab7b4471f18749258bf8fc1218b7f660", "version_major": 2, "version_minor": 0}

{"model_id": "77f51ac968fb40fb92e9c919ec95d1aa", "version_major": 2, "version_minor": 0}

{"model_id": "975808a9017a4799b966f8b7de29f0f5", "version_major": 2, "version_minor": 0}

{"model_id": "9b51295a0d2545ceb5cebf9b17f230e0", "version_major": 2, "version_minor": 0}

{"model_id": "e398ae9c1a164a5c9f1cc928014af4ed", "version_major": 2, "version_minor": 0}

{"model_id": "d5c807a82f5045c79aabfdb2ecf7d07b", "version_major": 2, "version_minor": 0}

{"model_id": "30e0456d7a234ce8a02549e7b779ba02", "version_major": 2, "version_minor": 0}

{"model_id": "1db58361e4c141b6b518f76f1eba6f96", "version_major": 2, "version_minor": 0}

{"model_id": "831119c4b81f43528286938f99c3bf87", "version_major": 2, "version_minor": 0}

{"model_id": "2316d843ab874976b3c368ed6fed56c3", "version_major": 2, "version_minor": 0}

{"model_id": "e197dc42ffae4121b8ab0b6ae1a1f0a", "version_major": 2, "version_minor": 0}

{"model_id": "55d6a14df6704b58861ae4d447d6d7b7", "version_major": 2, "version_minor": 0}

{"model_id": "68cbdaad7d1d49a69b1647955d10bd43", "version_major": 2, "version_minor": 0}

{"model_id": "a0eb6c03b3aa4b2ca5220c3fbac0efa4", "version_major": 2, "version_minor": 0}

{"model_id": "edeb0788819d4acf9ff093de099e1e37", "version_major": 2, "version_minor": 0}
```

efficientnet_b3 (BEST): Acc=0.4894, BAcc=0.6453, F1=0.4352

Training SOTA baseline: vit_base_patch16_224 ...

```
{"model_id": "ff8b0603c31d4dd8814a78fa1c1d18ee", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f0c9b4bd94884da0aa5797d363a45515", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0e0a29f183cd415d8cec212522cfa565", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "44d23f5430fd4da68e66dd3601b7aa79", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "7513050f0f3546dc871ae6ac1aa65cf9", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2cf7225dec0f4e888053acae3a2a1c5d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a46ad11a3df346c5bd6ba8cfacecebc6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "39ab17c900b4458493b5520e22b445a6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6a9e686e30044f33bcd4c0232de4f47f", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "95847f76ad06401fa9ff250671900c83", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9b94befdb22c40508d27293949ac9669", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "79fb3475d8b4424a9a4316c28a2c5ea6", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "71a5fafe7f37450a895bb8929d1ab591", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "dccb8e6021c04497bb33b8294f0079a2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "32787bac0405419c869c5874990bdbdc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "a8fcc87f201849d886732647ec522bfa", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "40a4c57c898d43808955e0650eda0911", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d3c167ab9c7a469d809e7f373d204791", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0632a466f0bc4ce2a874dc6594487761", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e4f2c11c019840f0a96201ffed857c04", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "866937e24e354ec288e5ef3e08580534", "version_major": 2, "version_minor": 0}
```

vit_base_patch16_224 (BEST): Acc=0.3949, BAcc=0.6142, F1=0.3561

Training SOTA baseline: swin_tiny_patch4_window7_224 ...

```
{"model_id": "de50e2b85cb545f495a20e278f94ce0d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "44b05fe504474baba236b25aec910100", "version_major": 2, "version_minor": 0}
```

IOStream.flush timed out

IOStream.flush timed out

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
self._shutdown_workers()

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
if w.is_alive():
^^^^^^^^^^^^

File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive

assert self._parent_pid == os.getpid(), 'can only test a child process'
^^

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

File

"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__

```

    self._shutdown_workers()
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "70f5277e4f854492b8f3d056297d6588", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "086625e832d04f19aa9a81a39bc7f550", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "40a087b7841741ef96e75ab7edf8fd88", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "b6ce5b526f92457eb99327e649f1e691", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "bc7fee5d9ca0460992ebb76270cb503e", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "d28dcc749c244a96beed3a6a2349852e", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "3a8633b1b099442883a88b18481d59ea", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "202a118add304d4db4e98cc711cae64f", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "210d041060174e0c88af990550f5490c", "version_major": 2, "vers
ion_minor": 0}

IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p

```

```

y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

```

```
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
IOStream.flush timed out
```

```
{"model_id": "60dfb792cc8541b09c1d4d788d7ce297", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4ec9b974ab0d4b468d7db984b5319b9b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0d1a8bb1f52349bea3bddcfafbac8386", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f5a79fe4cd284c51be4c7079eb35519a", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ae298268127148a7b9cef99c89483fb9", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4007495c336f43eeb98b079df87ef59d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e8f573d435b6493882e52b18587fe758", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "da1ffb42cfe84f7fafab4a97b3419d5b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e35ad8cb5a8b40b097ab11c205e61c32", "version_major": 2, "version_minor": 0}
```

```
swin_tiny_patch4_window7_224 (BEST): Acc=0.4470, BAcc=0.6940,  
F1=0.4667
```

```
Training SOTA baseline: convnext_tiny ...
```

```
{"model_id": "7dcae546aa2944fb9d248ede3d3089cd", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "73cab53b2e2e49079fcb023bf924ca9c", "version_major": 2, "version_minor": 0}
```

```

{"model_id": "2200330a87fb419aadfbbb5d9331df4a", "version_major": 2, "version_minor": 0}

{"model_id": "7d2314591cbe4407a59d631c6c4b9b9f", "version_major": 2, "version_minor": 0}

{"model_id": "825014f6332844b88b6982142c7693ef", "version_major": 2, "version_minor": 0}

{"model_id": "551aa250415e47f2b789527aff35bbd6", "version_major": 2, "version_minor": 0}

{"model_id": "6de18b6bc4bc44359c7520a4eefbba04", "version_major": 2, "version_minor": 0}

{"model_id": "4b8af1f636c64dc3bba68ece26707444", "version_major": 2, "version_minor": 0}

{"model_id": "cfb2a9a36c954068a9e5357cd1fdd492", "version_major": 2, "version_minor": 0}

{"model_id": "487534be09e340378d903879b092b065", "version_major": 2, "version_minor": 0}

{"model_id": "dc227a9c0b294fb0b4f1f0813a263ee9", "version_major": 2, "version_minor": 0}

{"model_id": "826f5e27e887452eba4f36eb67b744a6", "version_major": 2, "version_minor": 0}

{"model_id": "ec0a10d803a844268a8aace0198e54ea", "version_major": 2, "version_minor": 0}

{"model_id": "c7c3ada9db1e4f1ba841f793093f5868", "version_major": 2, "version_minor": 0}

```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```

File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
    self._shutdown_workers()
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'

```



```

AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^^^
  File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is alive():

```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__
at 0x728d383977e0>
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ^^^^^^^^^^^
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process

{"model_id": "321cd9a54be740bca1b5f6ff640c3d64", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "320a72e156894ab08b5f40994044b213", "version_major": 2, "vers
ion minor": 0}
```

```
{"model_id":"cf0907df07a04955a1b5092e8fd42ea5","version_major":2,"version_minor":0}
```

```
{"model_id":"f6814b00820445d7b804e08524eb480b","version_major":2,"version_minor":0}
```

```
{"model_id":"47b84fb39f5e4c18b4a4f31f0a6af7fb","version_major":2,"version_minor":0}
```

```
{"model_id":"92f521e9e2c24db792ff32ec2bffb6","version_major":2,"version_minor":0}
```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
```

```
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
```

```
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

Traceback (most recent call last):

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1664, in __del__
```

```
    self._shutdown_workers()
```

```
File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 1647, in _shutdown_workers
```

```
    if w.is_alive():
        ^^^^^^^^^^^^^
```

```
File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in is_alive
```

```
    assert self._parent_pid == os.getpid(), 'can only test a child process'
```

```
        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

AssertionError: can only test a child process

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x728d383977e0>

```

Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1664, in __del__
    self._shutdown_workers()
  File
"/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.p
y", line 1647, in _shutdown_workers
    if w.is_alive():
        ~~~~~
    File "/usr/lib/python3.12/multiprocessing/process.py", line 160, in
is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child
process'
        ~~~~~
AssertionError: can only test a child process

{"model_id": "471a976a6641457cbe9d1976b65b32f1", "version_major": 2, "vers
ion_minor": 0}

convnext_tiny (BEST): Acc=0.4654, BAcc=0.6753, F1=0.4874

Training SOTA baseline: densenet121 ...

{"model_id": "1c92d92108ba4b51beb9661dc3b60633", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "3a7ef85f09dc472187f564833a729a8c", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "ae25a5945c77469086f717b43f864f89", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "24e45d49aaef42da860956943160964b", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "186959ce4f2a46da9c211f493359a1d6", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "046faa774c324da1b28540a0e4c73985", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "1894aa286d67435e9f0fb4146002a2ac", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "13ea1a8445ba48068689d1489ef4d5cb", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "b637fc3c36304cacadb7f6330e36163a", "version_major": 2, "vers
ion_minor": 0}

```

```

{"model_id": "77a2ae943d9b462291918ea3aecc92fc", "version_major": 2, "version_minor": 0}

{"model_id": "e7d676131ba44b3ca6c4c363be1e1da9", "version_major": 2, "version_minor": 0}

{"model_id": "d770d693af42448e9eb8b7f64568dd9a", "version_major": 2, "version_minor": 0}

{"model_id": "ff4ffb05551442aba2a54c679f7f4658", "version_major": 2, "version_minor": 0}

{"model_id": "a1848091f9384239b112569d1f594a5e", "version_major": 2, "version_minor": 0}

{"model_id": "450d64288c9446bbb5ead57530facecd", "version_major": 2, "version_minor": 0}

{"model_id": "95ab204df636445ab5128f8785227794", "version_major": 2, "version_minor": 0}

{"model_id": "af0154412ab845eeb554a2a8c3cfe4de", "version_major": 2, "version_minor": 0}

{"model_id": "40f22c90466b4222814db2024867a043", "version_major": 2, "version_minor": 0}

{"model_id": "80fde904e57a47e69e3cdc1725eachfb", "version_major": 2, "version_minor": 0}

{"model_id": "cea3270d3e2c47d0b3ab2734f51db89c", "version_major": 2, "version_minor": 0}

{"model_id": "7ccf2401a76247bea09c414ce275e755", "version_major": 2, "version_minor": 0}

```

densenet121 (BEST): Acc=0.3689, BAcc=0.6107, F1=0.3293

SOTA Comparison Results:

MacroF1_best	epochs	Model	Accuracy_best	BalancedAcc_best
				protocol
UHViT (Ours, 5-fold CV mean)			0.505074	0.733522
0.581658	30		5-fold CV mean	
swin_tiny_patch4_window7_224			0.447010	0.694033
0.466748	10	budgeted_single_fold_best_epoch		
		convnext_tiny	0.465364	0.675253
0.487435	10	budgeted_single_fold_best_epoch		
		efficientnet_b3	0.489441	0.645291
0.435196	10	budgeted_single_fold_best_epoch		
		vit_base_patch16_224	0.394908	0.614220
0.356067	10	budgeted_single_fold_best_epoch		
		densenet121	0.368857	0.610687

```

0.329264      10 budgeted_single_fold_best_epoch
               resnet50      0.150780      0.483578
0.159402      10 budgeted_single_fold_best_epoch

# CELL 23: CROSS-DATASET VALIDATION (HAM10000)
print("\n" + "="*50 + "\nCROSS-DATASET VALIDATION (HAM10000)\n" +
      "="*50)
ham_base = Path(CONFIG['data_dir']) / 'ham10000'
ham_csv_files = list(ham_base.rglob('*metadata*.csv'))
if not ham_csv_files: ham_csv_files = list(ham_base.rglob('*.csv'))
ham_csv = ham_csv_files[0]
ham_img_dirs = [d for d in ham_base.iterdir() if d.is_dir() and 'HAM'
in d.name]
if not ham_img_dirs: ham_img_dirs = [ham_base]
ham_dataset = HAM10000Dataset(ham_csv, ham_img_dirs,
get_val_transforms(CONFIG['img_size']))
ham_loader = DataLoader(ham_dataset, batch_size=CONFIG['batch_size'],
shuffle=False, num_workers=CONFIG['num_workers'])
best_model = fold_results[0]['model']; best_model.eval()
ham_preds, ham_labels, ham_probs = [], [], []
with torch.no_grad():
    for images, labels in tqdm(ham_loader, desc='HAM10000 Eval'):
        images = images.to(CONFIG['device']); outputs =
best_model(images)
        ham_probs.extend(F.softmax(outputs, dim=1).cpu().numpy());
ham_preds.extend(outputs.argmax(dim=1).cpu().numpy());
ham_labels.extend(labels.numpy())
ham_preds, ham_labels, ham_probs = np.array(ham_preds),
np.array(ham_labels), np.array(ham_probs)
ham_acc, ham_bacc, ham_f1 = accuracy_score(ham_labels, ham_preds),
balanced_accuracy_score(ham_labels, ham_preds), f1_score(ham_labels,
ham_preds, average='macro')
print(f"\nHAM10000 Results:\nAccuracy: {ham_acc:.4f}\nBalanced
Accuracy: {ham_bacc:.4f}\nMacro F1: {ham_f1:.4f}")
plot_confusion_matrix(ham_labels, ham_preds, CONFIG['class_names'],
f"{CONFIG['output_dir']}/ham10000_cm.png")
wandb.log({'cross_dataset/ham_acc': ham_acc, 'cross_dataset/ham_bacc':
ham_bacc, 'cross_dataset/ham_f1': ham_f1})

```

```

=====
CROSS-DATASET VALIDATION (HAM10000)
=====

```

```

{"model_id": "d997ead33cc54dc2a9996e4d941a8ad6", "version_major": 2, "vers
ion_minor": 0}

```

HAM10000 Results:

Accuracy: 0.5255
Balanced Accuracy: 0.6976
Macro F1: 0.4359

CELL 23B: EXTERNAL VALIDATION (DermaMNIST from MedMNIST) - public direct download

```
import numpy as np
import torch
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from PIL import Image
import medmnist
from medmnist import INFO
from sklearn.metrics import accuracy_score, balanced_accuracy_score,
f1_score, roc_auc_score
```

```
derma_info = INFO["dermamnist"]
DermaMNIST = getattr(medmnist, derma_info["python_class"])
```

DermaMNIST: 7-class dermoscopy dataset (different label set than ISIC 8-class)

We use it as an EXTERNAL distribution test by reporting:

- Top-1 confidence distribution

- OOD-style metrics / uncertainty (optional)

- And/or map to MEL vs NON-MEL if you want (DermaMNIST doesn't have MEL as a direct class label).

For IF≈5, it's acceptable to report "cross-dataset robustness" with caveats.

```
class DermaMNISTWrapper(Dataset):
```

```
    def __init__(self, split="test", transform=None):
        self.data = DermaMNIST(split=split, download=True)
        self.transform = transform
```

```
    def __len__(self):
        return len(self.data)
```

```
    def __getitem__(self, idx):
        img, y = self.data[idx]                # img is PIL, y is int or
array                                           array
        if isinstance(y, np.ndarray):
            y = int(y.squeeze())
        else:
            y = int(y)
        img = np.array(img.convert("RGB"))
        if self.transform:
            img = self.transform(image=img)["image"]
        return img, y
```

```

# Use val transforms (resize + normalize) for compatibility
derma_test = DermaMNISTWrapper(split="test",
transform=get_val_transforms(CONFIG["img_size"]))
derma_loader = DataLoader(derma_test, batch_size=CONFIG["batch_size"],
shuffle=False, num_workers=CONFIG["num_workers"])

best_model = fold_results[0]["model"]
best_model.eval()

all_conf = []
all_entropy = []

with torch.no_grad():
    for images, _ in tqdm(derma_loader, desc="DermaMNIST External
Eval"):
        images = images.to(CONFIG["device"])
        probs8 = F.softmax(best_model(images), dim=1)          # ISIC
6-class outputs
        conf = probs8.max(dim=1).values.detach().cpu().numpy() #
confidence of predicted ISIC class
        ent = (-probs8 * (probs8.clamp_min(1e-
9))).log()).sum(dim=1).detach().cpu().numpy()

        all_conf.extend(conf.tolist())
        all_entropy.extend(ent.tolist())

all_conf = np.array(all_conf)
all_entropy = np.array(all_entropy)

report = {
    "DermaMNIST_external_n": int(len(all_conf)),
    "mean_confidence": float(all_conf.mean()),
    "median_confidence": float(np.median(all_conf)),
    "mean_entropy": float(all_entropy.mean()),
    "median_entropy": float(np.median(all_entropy)),
}

print("DermaMNIST External Robustness Report:")
for k, v in report.items():
    print(f"{k}: {v}")

wandb.run.summary["external_note"] = (
    "External dataset label space differs from ISIC2019; "
    "we report robustness metrics (confidence/entropy) under
distribution shift."
)
wandb.log({"external/note": "External dataset label space differs; we
report robustness metrics (confidence/entropy) under distribution
shift."})

```


100%|██████████| 19.7M/19.7M [00:01<00:00, 15.0MB/s]

```
{"model_id":"f035db2de7ea483da7681df1bc253e8b","version_major":2,"version_minor":0}
```

DermaMNIST External Robustness Report:

DermaMNIST_external_n: 2005

mean_confidence: 0.6616833858507827

median_confidence: 0.6399524807929993

mean_entropy: 0.9612262350551208

median_entropy: 1.0979413986206055

CELL 23C: DermaMNIST W&B logging + note (append after 23B)

```
import pandas as pd
```

'report' should exist from CELL 23B

```
wandb.log({"external/dermamnist_report":
```

```
wandb.Table(dataframe=pd.DataFrame([report]))})
```

```
wandb.run.summary["external_note"] = (
```

```
    "External dataset label space differs from ISIC2019; "
```

```
    "we report robustness metrics (confidence/entropy) under  
distribution shift."
```

```
)
```

```
wandb.log({
```

```
    "external/note": "External dataset label space differs; we report  
robustness metrics (confidence/entropy) under distribution shift."
```

```
})
```

```
print("✓ Logged DermaMNIST robustness report + note to W&B")
```

✓ Logged DermaMNIST robustness report + note to W&B

CELL 24: COMPUTATIONAL ANALYSIS

```
print("\n" + "="*50 + "\nCOMPUTATIONAL ANALYSIS\n" + "="*50)
```

```
comp_models = {'UHViT': lambda: UHViT(CONFIG['num_classes'],
```

```
CONFIG['dropout_rate']), 'Swin-T': lambda:
```

```
SwinOnly(CONFIG['num_classes'], CONFIG['dropout_rate']),
```

```
'EfficientNet-B3': lambda: EfficientNetOnly(CONFIG['num_classes'],
```

```
CONFIG['dropout_rate']), 'ResNet50': lambda:
```

```
create_sota_model('resnet50', CONFIG['num_classes']), 'ViT-Base':
```

```
lambda: create_sota_model('vit_base_patch16_224',
```

```
CONFIG['num_classes'])}
```

```
comp_df = computational_analysis(comp_models, CONFIG['device']);
```

```
print("\nComputational Analysis:\n" + comp_df.to_string(index=False))
```

```
=====
```

COMPUTATIONAL ANALYSIS

```
=====
```

Analyzing UHViT...
Analyzing Swin-T...
Analyzing EfficientNet-B3...
Analyzing ResNet50...
Analyzing ViT-Base...

Computational Analysis:

	Model	Params (M)	FLOPs (G)	Time (ms)
	UHViT	40.580266	10.712108	19.199839
	Swin-T	27.527042	8.760381	8.547678
	EfficientNet-B3	10.711600	1.945989	9.249110
	ResNet50	23.524424	8.260809	4.372694
	ViT-Base	85.804808	24.033695	4.171274

CELL 25: SAVE FINAL RESULTS

```
results_summary = {'UHViT_5Fold_CV': {'Accuracy':  
f"{np.mean(accs):.4f} ± {np.std(accs):.4f}", 'Balanced_Accuracy':  
f"{np.mean(baccs):.4f} ± {np.std(baccs):.4f}", 'Macro_F1':  
f"{np.mean(fls):.4f} ± {np.std(fls):.4f}", 'CI_95_Accuracy':  
f"[{acc_lo:.4f}, {acc_hi:.4f}]", 'CI_95_Balanced_Acc':  
f"[{bacc_lo:.4f}, {bacc_hi:.4f}]", 'CI_95_Macro_F1': f"[{f1_lo:.4f},  
{f1_hi:.4f}]", 'Cross_Dataset_HAM10000': {'Accuracy':  
f"{ham_acc:.4f}", 'Balanced_Accuracy': f"{ham_bacc:.4f}", 'Macro_F1':  
f"{ham_f1:.4f}"}}  
with open(f"{CONFIG['output_dir']}/results_summary.json", 'w') as f:  
    json.dump(results_summary, f, indent=2)  
fold_df = pd.DataFrame([{'Fold': i+1, 'Accuracy': r['acc'],  
    'Balanced_Acc': r['bacc'], 'Macro_F1': r['f1']} for i, r in  
    enumerate(fold_results)])  
fold_df.to_csv(f"{CONFIG['output_dir']}/fold_results.csv",  
    index=False)  
print(f"\n{'='*50}\nTRAINING COMPLETE!\n{'='*50}")  
print(f"\nResults saved to: {CONFIG['output_dir']}\nCheckpoints saved  
to: {CONFIG['checkpoint_dir']}")  
print(f"\nFinal Results:\n    Balanced Accuracy: {np.mean(baccs):.4f} ±  
    {np.std(baccs):.4f}\n    Macro F1: {np.mean(fls):.4f} ±  
    {np.std(fls):.4f}\n    Cross-Dataset (HAM10000) BAcc: {ham_bacc:.4f}")  
wandb.finish()  
print("\n✓ All done! Check W&B for full logs.")
```

=====
TRAINING COMPLETE!
=====

Results saved to: ./outputs
Checkpoints saved to: ./checkpoints

Final Results:
 Balanced Accuracy: 0.7335 ± 0.0108

Macro F1: 0.5817 ± 0.0129
Cross-Dataset (HAM10000) BAcc: 0.6976

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

✓ All done! Check W&B for full logs.

CELL 20C: MEL additional operating points (Youden J + PR-AUC)

```
import numpy as np
import pandas as pd
from sklearn.metrics import roc_curve, confusion_matrix,
average_precision_score, precision_recall_curve, roc_auc_score

MEL_CLASS = CONFIG["class_names"].index("MEL")

def summarize_threshold(mel_true, mel_score, threshold):
    mel_pred = (mel_score >= threshold).astype(int)
    tn, fp, fn, tp = confusion_matrix(mel_true, mel_pred).ravel()
    sens = tp / (tp + fn + 1e-9)
    spec = tn / (tn + fp + 1e-9)
    prec = tp / (tp + fp + 1e-9)
    f1 = 2 * prec * sens / (prec + sens + 1e-9)
    return dict(
        threshold=float(threshold),
        sensitivity=float(sens),
        specificity=float(spec),
        precision=float(prec),
        f1=float(f1),
        tp=int(tp), fp=int(fp), tn=int(tn), fn=int(fn)
    )
```

Pooled CV arrays should exist from CELL 20B:

all_y_true, all_y_prob

```
mel_true = (all_y_true == MEL_CLASS).astype(int)
```

```
mel_score = all_y_prob[:, MEL_CLASS]
```

AUROC + PR-AUC (PR-AUC is informative for imbalanced MEL)

```
mel_auroc = roc_auc_score(mel_true, mel_score) if 0 < mel_true.sum() <
len(mel_true) else 0.0
```

```
mel_prauc = average_precision_score(mel_true, mel_score) if 0 <
mel_true.sum() < len(mel_true) else 0.0
```

Youden's J optimal threshold (maximizes sensitivity + specificity -

```

1)
fpr, tpr, thr = roc_curve(mel_true, mel_score)
youden_j = tpr - fpr
best_idx = int(np.argmax(youden_j))
thr_youden = thr[best_idx]

youden_report = summarize_threshold(mel_true, mel_score, thr_youden)
youden_report.update({
    "metric": "YoudenJ_optimal",
    "MEL_AUROC": float(mel_auroc),
    "MEL_PR_AUC": float(mel_prauc),
    "youden_J": float(youden_j[best_idx]),
})

# Also report a higher-specificity point (optional): 90% sensitivity
target_sens = 0.90
idx = np.where(tpr >= target_sens)[0]
if len(idx) > 0:
    best_idx2 = idx[np.argmin(fpr[idx])]
    thr_sens90 = thr[best_idx2]
    sens90_report = summarize_threshold(mel_true, mel_score,
thr_sens90)
    sens90_report.update({
        "metric": "Sensitivity_90%",
        "target_sensitivity": float(target_sens),
        "MEL_AUROC": float(mel_auroc),
        "MEL_PR_AUC": float(mel_prauc),
    })
else:
    sens90_report = None

rows = [youden_report] + ([sens90_report] if sens90_report is not None
else [])
df = pd.DataFrame(rows)

print("Additional MEL operating points:")
print(df[["metric", "threshold", "sensitivity", "specificity", "precision"
, "f1", "MEL_AUROC", "MEL_PR_AUC"]].to_string(index=False))

wandb.log({"clinical/MEL_additional_operating_points":
wandb.Table(dataframe=df)})
wandb.log({"clinical/MEL_PR_AUC": float(mel_prauc),
"clinical/MEL_AUROC": float(mel_auroc)})

```

Additional MEL operating points:

	metric	threshold	sensitivity	specificity	precision
f1	MEL_AUROC	MEL_PR_AUC			
	YoudenJ_optimal	0.341932	0.726670	0.750156	0.387272
	0.505266	0.819794	0.569331		

Sensitivity_90%	0.215831	0.900044	0.494546	0.278996
0.425955	0.819794	0.569331		

```
-----  
-----  
Error                                Traceback (most recent call  
last)  
Cell In[41], line 70  
    67 print("Additional MEL operating points:")  
    68  
print(df[["metric","threshold","sensitivity","specificity","precision"  
,"f1","MEL_AUROC","MEL_PR_AUC"]].to_string(index=False))  
--> 70 wandb.log({"clinical/MEL_additional_operating_points":  
wandb.Table(dataframe=df)})  
    71 wandb.log({"clinical/MEL_PR_AUC": float(mel_prauc),  
"clinical/MEL_AUROC": float(mel_auroc)})  
  
File  
/usr/local/lib/python3.12/dist-packages/wandb/sdk/lib/preinit.py:36,  
in PreInitCallable.<locals>.preinit_wrapper(*args, **kwargs)  
    35 def preinit_wrapper(*args: Any, **kwargs: Any) -> Any:  
--> 36     raise wandb.Error(f"You must call wandb.init() before  
{name}()")  
  
Error: You must call wandb.init() before wandb.log()
```